

REVERSE ENGINEERING A BLUETOOTH APPLICATION: DISCOVERING THE SECRETS OF A MAKEUP PRINTING DEVICE

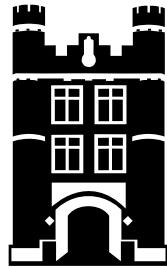
INDEPENDENT STUDY THESIS

Presented in Partial Fulfillment of the Requirements for
the Degree Bachelor of Arts in Computer Science in the
Department of Mathematical & Computational Sciences at
The College of Wooster

by
Natalie Pargas
The College of Wooster
2025

Advised by:

Professor Drew Guarnera (Computer
Science)



THE COLLEGE OF
WOOSTER

© 2025 by Natalie Pargas

ABSTRACT

The abundance and ubiquity of small, bespoke, Internet of Things (IoT) devices has widened a growing knowledge gap between what is understood about them and their actual capabilities. Thousands of new devices are released each year, where each variety can see millions of units produced and sold. Closing or narrowing this gap requires a coordinated analysis effort, necessitating bolstered numbers of individuals with the knowledge to address this issue. However, without access to source code, build system, and detailed software/hardware documentation, this cannot be done in the traditional analysis or forward engineering paradigm. In contrast, reverse engineering is the iterative process and methodology by which hardware and software can be understood with little or no access to the information and resources the developers used to make a given device. This paper serves as a guide on how to approach and successfully reverse engineer a device to the point of being able to augment its' functionality.

This work is dedicated to my mother.

ACKNOWLEDGMENTS

First, I would like to acknowledge my advisor, Professor Drew Guarnera. His guidance and encouragement during our advising sessions were instrumental in helping me see this project through. Despite the many challenges that arose along the way, his support consistently lit a path forward.

I also owe deep gratitude to my second advisor, Professor Heather Guarnera. I would not have even begun this project without her. She has championed me in every academic aspect, and it was her belief in me that allowed me to believe that this project was possible in the first place. I will forever be grateful for her counsel and motivation. She has gone far beyond what would be considered exceptional for an advisor, and I am so fortunate to have had her in my corner.

Next, I would like to thank Dean Bowen for her unwavering support with any college-related matters that threatened to impede my progress outside the classroom. She has been a reliable source of information and encouragement throughout this process. I have yet to meet someone as efficient and responsive when it comes to answering emails and finding solutions.

Finally, I would like to extend my gratitude to the entire Department of Mathematical and Computational Sciences. This has always been a department that truly cares about student learning and is willing to meet students where they are. It is through the hard work and dedication of my professors that I have been able to grow academically and, ultimately, complete this project.

CONTENTS

Abstract	iii
Dedication	iv
Acknowledgments	v
Contents	vi
List of Figures	ix
List of Tables	xi
List of Listings	xii
CHAPTER	PAGE
1 Introduction	1
1.1 What is Reverse Engineering?	1
2 Reverse Engineering Methods, Tools, and examples	5
2.0.1 Static Analysis	6
2.0.1.1 Disassembler	7
2.0.1.2 Decompiler	7
2.0.1.3 Dumping Tools	8
2.0.2 Dynamic Analysis	9
2.0.2.1 Debugger	9
2.0.2.1.1 User Mode Debuggers	10
2.0.2.1.2 Kernel Mode Debuggers	10
2.0.2.2 System Monitor	11
2.0.2.3 Virtual Machine	12
2.1 Protocol Analysis	13
2.2 Compiled Code RE	15
2.2.1 High level languages	15
2.2.1.1 Control Flow	16
2.2.2 Low level languages	17
2.2.2.1 Registers	18
2.2.2.2 The Stack	20
2.2.2.3 Flags	21
2.2.2.4 Functions	21
2.3 Example Problems	22

2.3.1	Reversing Hello World	23
2.3.2	Crack Me	25
3	Bluetooth	29
3.1	Bluetooth	29
3.1.1	Bluetooth Types	30
3.2	Pairing-Based Network	31
3.2.1	Piconet	31
3.2.2	Scatternet	33
3.3	BLE Mesh Networks	33
3.4	Broadcast Network from LE	34
3.5	Connections and Mixed Topology	35
3.6	SCO and ACL Logical Transports	36
3.6.1	BR/EDR Asynchronous Connection-Oriented (ACL)	36
3.6.2	BR/EDR Synchronous Connection-Oriented (SCO)	37
3.7	Bluetooth Low Energy Broadcasting and Observing	38
3.8	Bluetooth Protocols and Profiles	39
3.8.1	Generic Profiles	39
3.8.1.1	Generic Access Profile (GAP)	40
3.8.1.2	Generic Attribute Profile (GATT)	40
3.8.2	Use-Case-Specific Profiles	40
3.8.2.1	SIG-Defined GATT-Based Profiles	41
3.8.2.2	Vendor-Specific Profiles	41
3.9	Bluetooth Low Energy (BLE) Architecture	42
3.9.1	Application	43
3.9.2	Host	43
3.9.3	Controller	43
3.10	BLE Link Layer and Device Addressing	44
3.10.1	Link Layer States	44
3.10.2	Advertising, Scanning, and Connecting	46
3.10.3	Bluetooth Device Address	46
3.11	Host Controller Interface (HCI) Layer	47
3.11.1	Universal Asynchronous Rx/Tx (UART)	48
3.12	Connection	49
3.13	Services and Characteristics	50
3.13.1	Attribute Protocol (ATT)	50
3.13.1.1	Roles in ATT	50
3.13.1.2	Data Format: Attributes	51
3.14	Profiles	51
3.14.1	Generic Attribute Profile (GATT)	51
3.14.2	GATT Roles	52
3.14.3	Services	52
3.14.4	Characteristics	53
3.14.5	Profiles	53
3.15	Nordic Semiconductor	54

3.15.1	nRF Connect	54
4	Android Reverse Engineering	55
4.1	A Brief Overview of Android Fundamentals	57
4.1.1	Structure of an APK File	61
4.2	Overview of the YSL Makeup Printer	62
4.3	Reverse Engineering the Rouge Sur Mesure App	66
4.4	Environment Setup	67
4.4.1	Rooting the phone	67
4.4.2	Virtual Machine	70
4.4.2.1	JADX	71
4.4.2.2	Frida	71
4.4.2.3	Wireshark	72
4.5	Altering the app	72
5	Methodology	76
5.1	Decoding app data	76
5.1.1	Testing with the nRF Connect App	77
5.2	Testing with the nRF Sniffer and Wireshark	81
5.3	Analyzing the HCI Snoop Logs	83
5.4	Scripting	86
5.5	Analyzing BLE Payloads via Frida Hooking and Packet Inspection	88
5.6	Executing the Dispense	100
6	Conclusions and Future Work	106
6.1	Reverse Engineering Outcome	106
6.2	Future Work	108
	References	110

LIST OF FIGURES

Figure	Page
2.1 Protocol Analysis Process	15
2.2 Divide function code	18
2.3 Disassembly view in Ghidra	18
2.4 Decompiled code view in Ghidra	18
2.5 Register and Flag descriptions and purposes	20
2.6 Hello World Program	23
2.7 Hello World Disassembly	24
2.8 Crack Me Header	25
2.9 Crack Me Libraries	25
2.10 Crack Me Strings	26
2.11 Crack Me Ghidra Disassembly	27
2.12 Crack Me Code Section 1	27
2.13 Crack Me Code Section 2	28
2.14 Crack Me Code Executable	28
3.1 Bluetooth Devices	30
3.2 Central and Peripheral device relationship	31
3.3 Piconet Topology	32
3.4 Scatteret Topology	33
3.5 Mesh topology	34
3.6 Mixed Topology	36
3.7 Broadcasting and Observing	38
3.8 BLE Protocol Stack	42
3.9 Link Layer State Machine	45
3.10 Example Bluetooth Address	46
3.11 Host, HCI, and Controller components	48
4.1 Java Source to Hardware	56
4.2 Android Activity Layout	57
4.3 Android Activity to User	57
4.4 Androiod SDK Tools	58
4.5 Gradle Build Structure	59
4.6 Android App Components	62
4.7 Rouge Sur Mesure Custom Lip Color Creator	63

4.8	Color Wheel App Screen	64
4.9	Color from Image App Screen	64
4.10	Rouge Sur Mesure Color Families	65
4.11	Rouge Sur Mesure Cartridge Trios	65
4.12	Altered App Screen	75
5.1	nRF Connect Scanner page	77
5.2	nRF Connect Beam Beacon page	77
5.3	nRF Connect Beam Services and Characteristics	77
5.4	NRF Connect screenshots	77
5.5	Jadx App Gatt Profile Information	78
5.6	onConnectionStateChange() State Machine	79
5.7	nRF Connect Android Device Communication	80
5.8	Wireshark Output from Bluetooth Packet Sniffing	82
5.9	Wireshark Output from HCI Snoop Logs	84
5.10	HCI Packet Tree in Wireshark	85
5.11	Payload from Hooked Send Method	90
5.12	Beam Lips Device Model Jadx	94
5.13	Initial Payload from dispenseFor(...) Hook	94
5.14	Payload Two from dispenseFor(...) Hook	95
5.15	Device ID from the App	96
5.16	Payload Three from dispenseFor(...) Hook	96
5.17	Payload Four from dispenseFor(...) Hook	97
5.18	Payload Five from the dispenseFor(...) Hook	98
5.19	Payload Six from the dispenseFor(...) Hook	99
5.20	Multiple Payload Analysis	100
5.21	Dispense Environment Setup	101
5.22	Picture of Altered Dispense Output	104
5.23	Frida Console Logs During Altered Dispense	104
5.24	cn63 Screen Following Altered Dispense	105

LIST OF TABLES

Table	Page
5.1 Cartridge Metadata and Status at Time of Dispense	93

LIST OF LISTINGS

Listing	Page
---------	------

CHAPTER 1

INTRODUCTION

1.1 WHAT IS REVERSE ENGINEERING?

Reverse engineering is the process of disassembling, deconstructing, or analyzing a system, device, or piece of software in order to gain a deep understanding of how it functions. This concept has existed long before its modern applications in technology. For example, the dissections many of us performed in high school biology classes were, in essence, an early form of reverse engineering. Those dissections were carried out to break down something as complex and mystifying as life into understandable components and functions, revealing how various biological systems worked together. Similarly, one of the most effective ways to fully comprehend how a technological system operates is to open it up and inspect it piece by piece.

Reverse engineering software can be accomplished through a variety of techniques. Broadly speaking, these methods fall into two categories: static analysis and dynamic analysis. Static analysis involves studying a snapshot of the code in a single, unchanging state without executing it. This provides insight into how the software is structured, its logic, and potential areas of interest. In contrast, dynamic analysis takes place while the program is actively running, allowing an analyst to observe how it behaves in real time, including how it handles data and responds to inputs. Typically, a reverse engineering

process might begin by passing a program's executable to a disassembler, which translates machine code into assembly-level instructions. From there, a programmer might use system monitoring tools to gather information provided by the operating system about how the program interacts with its environment. Additionally, debuggers allow reversers to observe the CPU's internal operations, stepping through the disassembled code one instruction at a time to scrutinize the logic behind each operation. Using these various tools, a programmer must rely heavily on intuition, experience, and problem-solving skills to navigate the complex and often opaque process of reverse engineering.

There are numerous reasons someone might choose to reverse engineer software. The only time the inner workings of software are openly accessible to anyone is when it is released as open source. Open source software consists of freely available source code that anyone can read, modify, or distribute. However, developers frequently choose not to make their source code public for a range of reasons, including legal and licensing restrictions, protection of proprietary technologies, or personal business strategies. The opacity of proprietary software often leaves users, researchers, or competing developers with no choice but to reverse engineer if they wish to understand or interact with it.

Some of the most common motivations for reverse engineering software include performing malware analysis, improving one's programming skills, recovering lost source code, and enabling interoperability between different systems or applications. My own first introduction to the world of reverse engineering came through a video in which the creator was reverse engineering Apple's FaceTime for Mac. His goal was to reinstate closed captioning functionality for his hearing-impaired mother—a striking example of how reverse engineering can serve highly personal and practical purposes. The potential applications of reverse engineering are as diverse and far-reaching as the software programs that exist today.

Most software developers are familiar with reverse engineering primarily through its use in malware analysis. A classic example of malware is a Trojan virus, in which

malicious developers conceal harmful code inside programs that appear harmless. Reverse engineering such programs allows security researchers to break them apart and reveal the hidden malware lurking within. Additionally, modern software development often involves the use of external libraries. In large projects, keeping track of all dependencies and the interactions between different modules can be a formidable challenge. Hackers can exploit this complexity by hiding malicious entry points in these libraries to gain unauthorized access to sensitive information. In many cases, reverse engineering remains the only foolproof method for uncovering the true operations of a program and exposing any hidden threats or vulnerabilities embedded deep within the code.

Another compelling reason to reverse engineer software is to enhance one's own technical skills. The process requires extensive knowledge of both low-level assembly language and high-level programming concepts. Reverse engineering someone else's code is no simple feat. It forces the reverser to learn deeply about various methods of structuring and implementing software, as well as the diverse strategies developers use to solve problems. Through this process, reverse engineers often become far more proficient programmers, gaining insight into advanced coding techniques and architectural design choices that might not be visible in standard learning materials. This project itself will serve as an endeavor in learning not only forward engineering but also the intricate art and science of reverse engineering.

Another reason people may reverse engineer software is to recover lost source code. Though this is far from easy, it is sometimes possible to reconstruct high-level source code from compiled executables. Many developers have experienced the frustration and panic that comes with losing source code due to accidental deletion, hardware failures, or lack of proper backups. In such scenarios, reverse engineering can offer a potential solution, providing a lifeline to recover critical work that might otherwise be lost forever.

Interoperability represents yet another crucial motivation for reverse engineering. In many cases, achieving compatibility between different programs, systems, or devices is

impossible without reverse engineering. When working with external libraries or operating systems that provide documentation without accompanying source code, developers frequently encounter undocumented behaviors or missing details. While a programmer could attempt to resolve these issues through trial and error or by contacting the vendor, reverse engineering provides a definitive and often faster way to understand how to integrate different systems effectively. It serves as a bridge between incompatible software components, enabling seamless communication and cooperation in ways that might otherwise remain impossible.

This project serves as a case study in using reverse engineering to achieve interoperability with a proprietary Bluetooth device.

CHAPTER 2

REVERSE ENGINEERING METHODS, TOOLS, AND EXAMPLES

Because reverse engineering serves a wide variety of purposes, there is an equally broad range of methods and strategies employed to achieve those goals. In *Reversing: Secrets of Reverse Engineering*, these methods are generally categorized into two primary types. The first type focuses on large-scale observation of a program and is referred to as system-level reversing [7]. System-level reversing seeks to determine the overarching structure of a program and identify potential areas of interest within it. Once a general understanding of the program's layout has been established, the analysis can proceed to the second type: code-level reversing. Code-level reversing delves deeper into the specifics of a program's inner workings, offering detailed insights into particular segments of code and uncovering how algorithms and functionalities are implemented.

System-level reversing is achieved by running an array of specialized tools on a program and observing the services provided by the operating system. These observations gather information on how the program interacts with system resources, tracks input and output operations, and reveals characteristics of the program's executable files. Much of this data is supplied by the operating system, as any portion of a program that interacts with components outside itself must necessarily communicate through the OS. As a result, effective system-level reversing demands a strong familiarity with the operating system's internal workings, making it invaluable for reverse engineers to develop deep expertise in the OS environment they are targeting.

Code-level reversing, by contrast, is a far more complex and meticulous process. It

involves extracting algorithms, logic, and design concepts directly from a program's binary code. This task requires a robust background in reverse engineering techniques, software development practices, CPU architecture, and operating system internals. Even when working with readable, well-documented source code, modern software systems can be enormously complex. For instance, the program examined in this project comprises 21 frameworks, contributions from at least three separate companies, and potentially tens or even hundreds of thousands of lines of code. Moreover, when high-level code is compiled into binary instructions, a single line of code can transform into ten or more lines of low-level machine instructions. Reversing production-level software from binary back to source code is thus an almost incomprehensibly difficult endeavor. Thankfully, multiple methodologies and a variety of specialized tools have been developed to aid in effective code-level reversing, making this seemingly impossible task more approachable for skilled practitioners.

2.0.1 STATIC ANALYSIS

Static analysis is the practice of analyzing code without executing it. This approach allows reverse engineers to study the structural and logical components of a program in a controlled environment, free from the risks associated with running potentially malicious code. Static analysis can also serve as a valuable precursor to dynamic analysis, helping to identify key areas of interest for more in-depth examination later.

Static analysis tools operate at both the source code level, when code has not yet been compiled, and at the machine language level, analyzing binary or assembly-level code. Reverse engineers frequently use static analysis tools to disassemble machine code into assembly instructions, enabling them to examine and interpret how different sections of code relate to one another. While retrieving disassembled code is one of the most fundamental functions of static analysis tools, many modern tools offer a host of additional features, such as graphical representations of control flow, cross-referencing functions, and

identifying potential vulnerabilities. These tools are often designed to provide reverse engineers with a starting point, helping them navigate the often bewildering world of assembly-level functions and binary logic.

2.0.1.1 DISASSEMBLER

One of the most key static analysis tools for a reverse engineer is a disassembler. Disassemblers are applications that take in a program's executable file and generate a file with the assembly code. This isn't too difficult because assembly is just a text translation of the machine readable binary code. Unfortunately, that makes it a lot less clear than higher level programming languages. Disassembly is also unique to a person's processor, but there are disassemblers that can support more than one CPU architecture. Some examples of disassemblers are IDA, Binary Ninja, Hopper, Ollydbug, Radare2, and Ghidra.

2.0.1.2 DECOMPILER

Decompilers are similar but more complex than disassemblers. Instead of just producing the assembly language code, decompilers attempt to produce high-level readable code. They attempt to produce something that looks as close to the source code as possible by reversing the compilation process [7]. The front end of the decompiler decodes low-level assembly instructions and translates it into an intermediate representation specific to the decompiler. The intermediate representation is then iterated on to remove as many extraneous details as possible while preserving the important details. Lastly the back end takes the polished up intermediate representation and translates it again into a high level language.

While this may sound like a simple way to retrieve the source code, it is highly unlikely the high level language returned will be something that actually matches. Anyone who has ever run text through a translation website multiple times knows that each iteration causes a loss in fidelity and oftentimes the end result will be gibberish. Decompilers go a

step further and have no immediate knowledge of things like function or variable names and structures. This does not make them useless, though. For example, a text translator will usually return your original result when given a simple phrase. Similarly a decompiler will most likely be able to pick up on more straightforward functions such as adding a + b. Decompilers offer hints and snippets into what the source code may look like which is valuable information for a reverse engineer. Most of the examples for disassemblers listed also count as decompilers.

2.0.1.3 DUMPING TOOLS

Executable dumping is usually the first step in reverse engineering a program. Executable dumping helps inform a reverse engineer what a program does, how it interacts with external elements, and general knowledge of how a file is structured. Using dumping tools, a reverse engineer will start by figuring out what type of file they're dealing with. They will then start to unpack the format. Extracting strings will immediately offer insight into what language the source code is written in, what library modules it uses, what kinds of data it is dealing with, and what the goals of program functions may be. Other useful information we can retrieve is the layout of the file in memory and what the general logical structure of the file is. Once we know the type of file we also know which tools we can use. For example, some debuggers only work with certain high and low level languages. Some popular executable dumping programs are ones already built into your machine like DUMPBIN for Windows or otool for Mac.

Otool is identifies a file's Mach header. On systems with a Mach kernel like macOS and iOS, the executable binaries they use will have a Mach header. All headers have a magic number to identify them. Thin binaries only have that number while fat binaries with fat headers will have locations of the other executable's headers in them.

2.0.2 DYNAMIC ANALYSIS

Dynamic analysis occurs when a program—or specific sections of it—is executed during the analysis process. This approach provides unique insights into how software behaves in real time, revealing execution paths, input/output relationships, and interactions with external systems. Because running an unknown program can produce unpredictable and potentially harmful results, it is safest to conduct dynamic analysis within a secure, isolated, or sandboxed environment. Virtual machines are commonly used for this purpose, as they are easy to set up, control, and reset between tests.

A wide variety of tools contribute to dynamic analysis. Any tool that executes code and monitors its behavior falls under this category. Some dynamic analysis tools run entire programs to observe overall functionality, while others focus on specific sections, functions, or even single instructions. Additionally, tools that monitor a program's external environment during execution—for example, capturing network traffic, tracking file system changes, or observing registry modifications—are also considered part of dynamic analysis. Together, these tools give reverse engineers a powerful way to understand not just how a program is built, but how it behaves in practice.

2.0.2.1 DEBUGGER

Debuggers are a tool usually used for developers to locate and work through errors in their programs, but they're also vital tools in reverse engineering. Many debuggers can work through assembly language. While assembly language may be difficult for a human to parse through, it is the exact same logic that gets broken down and sent to a computer meaning that the computer can show what each assembly instruction is intended to do. Most debuggers software engineers use were actually designed from the ground up with the purpose of stepping through assembly code. The debugger can show the state of CPU registers along with a memory dump that shows what's in the stack. Debuggers usually contain disassemblers which as discussed is a vital reverse engineering tool. Many

debuggers also contain both software and hardware breakpoints. Software breakpoints are instructions added into the code during runtime to pause and hand control over to the debugger. Hardware breakpoints are a CPU feature which allow the processor to pause execution when a certain memory address is reached and hand control over to the debugger. A reverse engineer can use insert breakpoints at data structures of interest and use the debugger to reveal what they are. Debuggers also offer a clear view of registers and memory to aid in a reverse engineer's ability to process low level information.

2.0.2.1.1 USER MODE DEBUGGERS Most debuggers that a programmer will use are user mode debuggers. They run on a system like any other application then seize control of the target program to debug [7]. An advantage to them is that they are easier to set up and navigate than their kernel-mode counterparts. Usually it's fine to stay limited to user mode viewing of an application. It only creates troublesome limitations when the target application has kernel mode components such as device drivers. User mode debuggers are also not always sufficient when trying to debug a program before it reaches the main entry point. These kinds of programs are usually ones that have a lot of statically linked libraries in the executable. The final and most likely to be problematic issue is that user-mode debuggers can only view a single process in a program. This can cause issues when the target application has processes that interact in unknown ways. The user may not know which process to zero in on that has the code of interest.

2.0.2.1.2 KERNEL MODE DEBUGGERS Kernel mode debuggers are different from user-mode debuggers because they capture the entire system and not just a single process. Instead of running atop the operating system, kernel mode debuggers sit alongside the system kernel and stay ready to capture the entire system's stats. Kernel mode debuggers can be more helpful to reverse engineers because they offer more clues as to what's going on with the system. A key tool for kernel-mode debuggers is that they allow the placement of low level code breakpoints. As a reverse engineer working with assembly code, being

able to test different sets of assembly instructions that may be the code section a reverser is looking for is incredibly useful.

For example, picture a scenario where a reverse engineer is looking for the API responsible for handling moving windows? That will need to be managed by a windows manager in the kernel. The complexity is introduced when trying to identify which API moves a particular window. Since there are multiple APIs that can be used to move a window, pinpointing the one you are trying to focus on is a difficult task. This is where kernel mode debuggers come in handy. Low level breakpoints in the operating system responsible for shifting windows around will help you identify which API is getting called when a window is moved [7]. From there the reverse engineer will be able to hone in on that API.

Unfortunately, kernel mode debuggers come with their own drawbacks. Setting them up can be challenging because they require access to the full system. They need to suspend the entire system while running so they can go line by line, which means the system can no longer have multiple threads going [7]. They will also usually need to be set up inside a virtual machine which will be discussed in a later section [4]. These are reasons why a reverser should exhaust their other options before jumping into using a kernel mode debugger. Still, they are powerful tools when the scenario calls for it.

2.0.2.2 SYSTEM MONITOR

System monitoring can be a crucial part of the reverse engineering process. In some cases, a reverse engineer can figure out what they need through system monitoring tools without ever looking at the decompiled code. System monitoring tools can capture what happens between the code and hardware using the intermediate channels of input/output [7]. For local system monitoring, tools monitor things such as file operations (creating, deleting, moving, etc). For programs that communicate over networks, system monitoring might

look like recording all TCP/UDP network traffic. System monitoring tools are commonly used in virtual machines.

2.0.2.3 VIRTUAL MACHINE

Many programmers will encounter virtual machines (VM) at some point during their career. It's not uncommon for an application to interface exclusively with one type of operating system. One reason is because writing applications to work with different operating systems adds time and complexity that not all programmers can afford. Choosing to write in a high level language may make code more portable because there's more verbose built-in libraries or interpreters that deal with the specifics of the systems without the need for any programmer intervention. But high level languages often trade performance for convenience which is not a trade that a developer working with large amounts of data can necessarily afford to take. What happens when a developer has an application that runs on Windows while they have a Mac? This is where a virtual machine enters into the mix. Virtual machines are safe sandboxed environments that work as miniature computers with their own operating systems within a computer. They are constructed with an interpreter and bytecode. During compile time, select code snippets are compiled specific for the VM target architecture and then inserted into the program along with the interpreter. During run-time, the interpreter begins executing the bytecode. Virtual machines are costly to implement because they are so expansive which is why only the necessary code snippets get rendered [16]. Because virtual machines work in their own isolated environment, they offer powerful protections against unknown or potentially malicious software. This makes them a useful tool for reverse engineers who deal with software that they do not know the contents of. The level of control over virtual machines also makes them a useful tool because there is ultimate knowledge of system conditions.

2.1 PROTOCOL ANALYSIS

Communication protocols form the essential foundation for structured exchanges of information between different systems and components. Many applications, especially those that involve interoperability between different platforms or devices, depend on well-defined formats for sending and receiving messages. However, in cases where official protocol specifications are unavailable, incomplete, or proprietary, protocol reverse engineering can be a crucial tool for discovering the format, flow, and logic behind these communications.

Protocols are extensively used in telecommunications, networking, and distributed systems because they establish consistent rules for exchanging data between different entities. Protocols may be published as open standards, freely accessible for anyone to implement, or they might be kept proprietary, obscuring their details from end users and competitors. Protocol reverse engineering is most often employed in the latter scenario, where understanding the communication mechanisms is critical but official documentation is unavailable. The primary objective of protocol reverse engineering is to derive a model of how the target protocol functions, without any prior knowledge of its inner workings.

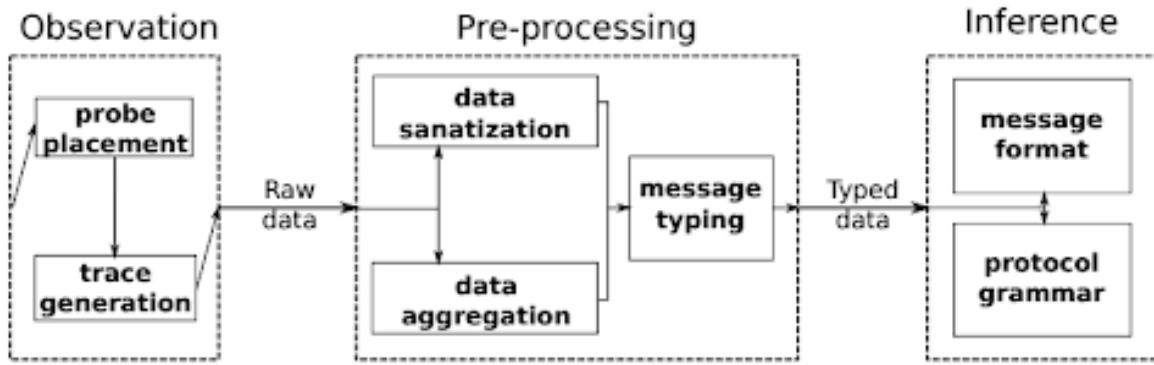
Protocol reverse engineering also plays a vital role in simulating network protocols. Network simulation allows researchers and developers to rapidly prototype specific test cases that might be cumbersome, time-consuming, or even impossible to execute in real-world environments. Furthermore, simulation tools can replay captured network traces under different conditions, helping engineers adapt their systems to new requirements or environments. In cybersecurity, protocol reverse engineering is frequently used in software security audits to verify that components behave correctly and securely under a range of circumstances, identifying potential vulnerabilities or compliance issues before they can be exploited.

Another significant application of protocol reverse engineering is network conformance testing, where the protocol in question is already known. In these cases, reverse engineering

can help create a detailed model of the protocol's behavior, which can then be compared to official standards or specifications to ensure that an implementation adheres to the expected rules and practices.

There are two general categories of protocol reverse engineering techniques. The first focuses on analyzing the messages exchanged between two components on a network. By carefully examining these exchanges—often referred to as network traces—analysts can infer the structure, commands, and responses that define the protocol's behavior. The second approach involves reverse engineering the application itself. This can mean analyzing the source code (when available), disassembling binary code, or tracing sequences of binary instructions during execution to reveal how the protocol logic is implemented internally [6].

Many tools exist to automate significant parts of the protocol reverse engineering process. The initial setup phase typically involves identifying and characterizing the operating environment in which the protocol operates. With a solid understanding of the environment, analysts can begin the observation phase, which involves using specialized tools to collect network traffic or execution traces. Once data has been gathered, it must be sanitized to extract only the relevant messages associated with the target protocol, filtering out noise and unrelated data. The final step is the actual analysis, where analysts draw inferences about the protocol's structure, message formats, and state machines. These steps are highly iterative and rely heavily on the expertise and insight of the analysts, as well as the capabilities of the tools being employed. Figure 2.1 illustrates this step-by-step process.

Figure 2.1: Protocol Analysis Process

2.2 COMPILED CODE RE

It is important for a reverser to understand the different goals of high level and low level languages. Understanding these differences will help with interpretation of why code is structured a certain way.

2.2.1 HIGH LEVEL LANGUAGES

High level languages exist as an abstraction to interface with operating system software and underlying hardware. High level languages work so a programmer can focus on specifying program implementation logic and structural design without needing to dedicate time to figuring out system specific details. If a programmer were to decide to write instructions to the machine they would use assembly language, a low-level, human readable language used to communicate more directly with the computer hardware. Those instructions would then get turned into machine code, an even lower level programming language that is readable to a computer's processor. Compiled languages are very performant as they are converted to assembly language then machine code. Writing directly in assembly language could allow for optimizations, but compilers are pretty good at enacting those

optimizations themselves. When interpreters or VMs (like the Java Virtual Machine) are involved, the VM invokes the necessary behaviors unlike in languages such as C or C++ where the compiled code becomes machine code for the target platform to run directly. While writing in something like assembly language can be very performant, it would be significantly more difficult and time consuming for a person to write a modern application in assembly. High level languages remove the need for programmers to understand their program at the assembly level, a reverse engineer however is forced to examine the low-level assembly version of a compiled program to understand the specific ways a program's goals are being achieved. The main goal of a reverse engineer is to use what they know about a program's behavior, the potential ways that can be accomplished in high level code, and how to read assembly code to determine specific instrumentation details. For reverse engineers, the most important thing to know about a high level language is to what level does it abstract or conceal the underlying machine code [7].

Languages like Python that have a built in interpreter will abstract it a lot. These programs may be full of extraneous instructions to perform necessary operations. On the other hand, languages like C are written much closer to the target processor and won't have nearly as much separation between the source and machine code

2.2.1.1 CONTROL FLOW

Control flow is what makes code determine what instructions get executed and how often. This manifests through statements like conditionals that give general instructions for what a program should do when a condition is true or false. A processor has no knowledge what statements like 'if' or 'while' mean. Under the hood, these statements translate into verbose and daunting assembly code. This is because high level conditional statements are often broken down into a series of smaller operation sequences to match the instruction set the processor understands [7].

Other typical structural components of programs that aid in control flow are switch

blocks and loops. Switch blocks, or $n - way$ conditionals, take in an input and have n number of code blocks that are potentially executed depending on the input value. Each block of code gets assigned at least one value prior to runtime. The compiler also generates code to receive the input value and search for the proper code block to execute. The values for the code blocks are usually stored in a lookup table that has pointers to each corresponding code block. Depending on the input value, the program will go through the process of searching the lookup table then jumping to the proper code block at runtime [7]. Loops work to allow a program to repeatedly execute a certain code block multiple times. A loop has a counter to keep track of how many iterations it has performed. There is also a conditional statement that determines when the loop will stop. Loops and conditionals are inherently intertwined. The difference is that loops execute over and over until the condition is no longer met.

To understand control flow sequences, one must understand how low level control flow is implemented. This means that a reverse engineer needs to know the specific rules of each kind of low level architecture because low level control flow instructions are unique to the platform for which they are designed.

2.2.2 LOW LEVEL LANGUAGES

Earlier we mentioned that complexity is introduced when working with low level system specific details. This is especially true when the goal is to translate high level logic into something that will be understood by your machine. Of course, there must be a way to do this because the CPU does it anytime a modern application is run. But it often involves keeping track of more details than a sole human is capable of. That is why a reverse engineer must develop the skill of parsing through assembly code and being able to create a kind of ‘mental image’ of how it relates to high level constructs. Data management is one of the things that a person’s computer keeps track of that would be difficult for a human to

do. To understand why this is, we can look at the data management of a relatively low level language, C, versus the assembly representation. Consider the code:

Figure 2.2: Divide function code

```
int divide(int a, int b)
{
    int result;
    result = a / b;
    return result;
}
```

Figure 2.3: Disassembly view in Ghidra

```
int divide(int a,int b)
{
    int iVar1;
    int in_w0;
    int in_w1;
    int result;

    iVar1 = 0;
    if (in_w1 != 0) {
        iVar1 = in_w0 / in_w1;
    }
    return iVar1;
}
```

Figure 2.4: Decompiled code view in Ghidra

While this function may seem incredibly simple, there is no direct translation into a machine code representation. To execute this in a low level language, it would require first storing the machine state. Then memory would need to get allocated for result. Variables a and b would need to get loaded from memory into a register. Then a would need to be divided by b and the result would need to get stored in the register that got loaded in the beginning. The machine state from before would need to get reloaded. The pointer would need to return to the caller and bring back result [7]. One line of high level code could result in any number of assembly instructions. Managing data is one of the biggest challenges of reverse engineering.

2.2.2.1 REGISTERS

To keep from needing to constantly access RAM, microprocessors have a smaller internal memory that can be accessed with barely any performance cost [7]. There are multiple types of internal memories that microprocessors keep and registers are one of them. Registers are little pieces of internal memory that live within the processor and are able to be accessed with an imperceptible cost to performance [4]. Data must be in a register in order for the CPU to access it for processing. The biggest problem with registers is that there are not

very many of them. One of the most popular processors, the IA-32, only has eight 32-bit registers that can be used for anything. It does have more registers, but they all have highly specific use cases. Assembly language is written around utilizing registers because they're so performant. They are not good for long term storage, though, that is when RAM becomes the better choice.

RAM is volatile memory, meaning its contents are lost when the computer is powered off. It temporarily holds data and instructions for currently active tasks, giving the CPU quick access to the broader context of running programs. The CPU operates using something called the instruction cycle: first, it fetches an instruction; then, it decodes or interprets it to determine the required operation; finally, it executes the instruction, carrying out the specified task.

One of the most important things to take away is that CPUs do not automatically manage these tasks. Data management is outlined in the assembly code which is what reverse engineers will need to get comfortable sorting through. One thing that reverse engineers will need to do is focus on figuring out what kind of values are getting loaded into a register. For example, it is easy to see when a register is only being used to grant instructions access to a specific value because that register will only show up when transferring value from memory to the instruction or vice versa. Another example is when a register shows up many times in one function. This is a good clue that one of the function's local variables is being stored in that register. Figure 2.5 shows some of the different registers and their purposes.

Figure 2.5: Register and Flag descriptions and purposes

General-purpose registers		Segment registers		
31	0	15	0	EIP
		EAX	code segment register	The instruction pointer (EIP) register contains the offset in the current code segment for the next instruction to be executed. It is advanced from one instruction boundary to the next in straight-line code or it is moved ahead or backwards by a number of instructions when executing JMP, Jcc, CALL, RET, and IRET instructions.
		EBX	data segment register	
		ECX	stack segment register	
		EDX	data segment register	
		ESI	data segment register	
		EDI	data segment register	
		EBP		
		ESP		
Status and control registers		EFLAGS		EFLAGS
31	0	EIP		The 32-bit EFLAGS register contains a group of status flags, a control flag, and a group of system flags. Following initialization of the processor (either by asserting the RESET pin or the INIT pin), the state of the EFLAGS register is 00000002H. Bits 1, 3, 5, 15, and 22 through 31 of this register are reserved. Software should not use or depend on the states of any of these bits. The following instructions can be used to move groups of flags to and from the procedure stack or the EAX register: LAHF, SAHF, PUSHF, PUSHFD, POPF, and POPFD. After the contents of the EFLAGS register have been transferred to the procedure stack or EAX register, the flags can be examined and modified using the processor's bit manipulation instructions (BT, BTS, BTR, and BTC)

1

2.2.2.2 THE STACK

The stack is one of two places that a value can be set aside. Registers are not managed by a processor and to use one you just need to load a value in it. Often there will not be registers available or there is a particular reason a variable will need to reside in RAM instead of in a register. That's when you'd put a variable on the stack instead [7].

The stack is a short term storage space in memory that gets utilized by both the CPU and the program. It is the spot where short term information gets put when a register can not be used for whatever reason. Registers are for the shortest term data while the stack is for the second shortest. The stack lives in RAM like all other data and is just a carved out section for intermediate term data. Modern operating systems tend to manage multiple stacks at the same time. Each of the concurrently running stacks is a representation of a program or thread [16].

¹Figure modified using Intel® 64 and IA-32 Architectures Software Developer's Manual and <https://flint.cs.yale.edu/cs422/doc/pc-arch.html#register>

Stacks use Last In First Out data management. Items are pushed on the top of the stack and popped from the top of the stack. Memory in stacks is allocated from the top down where the first in address is allocated and used first while the stack grows backwards towards lower addresses [7].

2.2.2.3 FLAGS

One of the registers you may have noticed from the previous figure 2.5 is the EFLAGS register. This is a collection of special IA-32 registers that contain system and status flags. The system flags' job is to manage the different modes and states of the processor. The status flags are what a reverse engineer will typically be more interested in and are used by the processor to record its current logical state [7]. They are often updated by various logical and integer instructions so they can record the outcome of these actions. There are also instructions whose operation conditions are dependent on the values for the status flags. This is what allows sequences of instructions to run different operations depending on what different input values may be, etc.

Flags in IA-32 are the crux of conditional code. Arithmetic instructions check operand conditions and then set processor flag conditions based on the resulting values. There are also sets of instructions designed to read the flags and do different operations based on the value. An example of a popular instruction set is Jcc or the Conditional Jump. It tests for predefined flag values and then jump depending on whether or not it matches with the specified conditional code [10].

2.2.2.4 FUNCTIONS

Instructions are the actual actions specified in assembly code. They are formatted with operation code (opcode) and one or two operands [7] The opcode is what you are asking the computer to do such as MOV or JMP. The operand is the parameters that get passed to the opcode or the data being manipulated. Certain instructions will have no parameters.

Data in assembly comes in three basic forms. There are register names, immediate data, and memory addresses. Register names are the names of the general purpose registers to be read from or written to like the EAX, EBX, etc. Immediate data is constant values that are embedded into the code and usually indicates there was a hard coded value in the source code. Memory addresses are the locations of operands stored in RAM. It can be hard coded and tell the processor where to read to and write from. It could also be a register with a value that will be used as a memory address. It is also possible to combine a register with an arithmetic operation and a constant so that there is some base address and then an index offset [7].

General purpose instructions are the ones that deal with program flow, logic, arithmetic, string operations, and basic data movement. They deal with data stored in memory at the general purpose registers, EFLAGS register, segment registers, and address information stored in memory [7].

The MOV instruction shows up most frequently in most IA-32 instruction sets. This one deals with basic data movement. It takes in a destination operand and a source operand then moves the data from the source to the destination. The sources can be registers, immediate, or memory addresses. It is important to note that MOV cannot transfer data through memory, it can only take it out or put it in. The destination address can be a memory address (using a register or an immediate) or a register [7].

2.3 EXAMPLE PROBLEMS

With some grounding in reverse engineering tools, methods, and assembly, it's a good time to try out a few practice problems. There are a couple of ways to approach this. One option is to work on “CrackMe” challenges, which are small programs designed to test reverse engineering skills by making analysis deliberately difficult. Another approach is to

reverse a self-written program, placing markers in the code to check whether they can be identified during analysis.

2.3.1 REVERSING HELLO WORLD

Before taking on the larger endeavor of reversing a program written in a declarative language (SwiftUI), it's important to start with the basics. Figure 2.6 shows a simple 'Hello World' program in C. To add some complexity, a few random variables and data structures were thrown in as Easter eggs. This is the program whose decompilation will be used as a jumping off point. Figure 2.7 is the decompilation of the program once it

Figure 2.6: Hello World Program

```
test4.cpp > main()
1 #include <iostream>
2
3 struct dope{
4     char b;
5     unsigned long long val;
6     short val2;
7 };
8
9 int main() {
10     int array[10] = {0};
11     array[8] = 1;
12     struct dope local_dope;
13     local_dope.b = 'A';
14     local_dope.val = 0xffffffff;
15     local_dope.val2 = 32000;
16
17     std::cout << "Hello World!";
18     return 0;
19 }
```

was put through Ghidra. The first thing to note is how with even a relatively low level language like C, the instructions in the code more than doubled. The disassembly being shown in the figure is also only a portion of Ghidra's output. The figure only shows the function section of the program output when imports, exports, labels, classes, etc take up the large majority of the program. One important part of reverse engineering is sorting

through all of the information to find out what exactly is important to the reverser. Some reverse engineers refer to this process as finding the shape and edges of the data. With the

Figure 2.7: Hello World Disassembly

Listing: a.out

```
***** undefined entry() *****  
undefined     w0:           <RETURN>  
undefined8    Stack[-0x10]:8 local_10  
  
undefined8    Stack[-0x18]:8 local_18  
  
undefined4    Stack[-0x20]:4 local_20  
undefined4    Stack[-0x44]:4 local_44  
undefined2    Stack[-0x50]:2 local_50  
undefined8    Stack[-0x58]:8 local_58  
undefined8    Stack[-0x60]:1 local_60  
  
_main          entry  
***** XREF[2]: Entry Point(*), 100003368(*) *****  
000003124 ff 83 01 d1      sub    sp,sp,#0x60  
000003128 fd 7b 65 a9      stp    x29,x30,[sp, #local_10]  
00000312c fd 43 01 91      add    x29,sp,#0x50  
000003130 08 00 00 b0      adrp   x8,0x100004000  
000003134 08 49 40 f9      ldr    x8,[x8, #offset ->__stack_chk_guard] = 10000c018  
000003138 08 01 40 f9      ldr    x8,[x8]=__stack_chk_guard = ??  
00000313c a8 83 01 f8      stur   x8,[x29, #local_18]  
000003140 01 #00 88 52      mov    w1,#0x0  
000003144 ff 1f 00 b9      str    wzr,[sp, #local_44]  
000003148 e8 83 00 91      add    x0,sp,#0x20  
00000314c 02 05 00 d2      mov    x2,#0x28  
000003150 37 03 00 94      bl    _memset  
000003154 28 00 88 52      mov    w8,#0x1  
000003158 e8 43 00 b9      str    w8,[sp, #local_20]  
  
#0x1 = A in hexadecimal  
local_20 is where 'A' is being stored  
w8 is address of local_20  
  
00000315c 28 08 80 52      mov    w8,#0x41  
000003160 e8 03 00 39      strb  w8,[sp]=local_50  
000003164 e8 7f 4b 02      orr    x8,xzr,#0xffffffff  
000003168 e8 07 00 f9      str    x8,[sp, #local_58]  
00000316c 08 a0 8f 52      mov    w8,#0x7d00  
000003170 e8 23 00 79      strh  w8,[sp, #local_50]  
000003174 00 00 00 b0      adrp   x0,0x100004000  
000003178 00 1c 40 f9      ldr    x0=>_ZN5t3__14coutE,[x0, #offset ->__ZN5t3__14coutE] = ??  
= 10000c000  
  
00000317c 01 00 00 90      adrp   x1,0x100003000  
000003180 21 00 2b 91      add    x1=>_Hello_World!_100003ef0,x1,#0xef0 = "Hello World!"  
000003184 0e 00 00 94      bl    __ZN5t3__1lsB8ue1700061NS11char_traitsIcEEERNS_13basic_ostrreamI... undefined __ZN5t3__1lsB8ue1700061NS11char  
000003188 a9 83 5f f8      ldur   x9,[x29, #local_18]  
00000318c 08 00 00 b0      adrp   x8,0x100004000  
000003190 08 49 40 f9      ldr    x8,[x8, #offset ->__stack_chk_guard] = 10000c018  
000003194 08 01 40 f9      ldr    x8,[x8]=__stack_chk_guard = ??  
000003198 08 01 09 eb      subs   x8,x8,x9  
00000319c e8 17 0f 1a      cset   w8,eq  
0000031a0 68 00 00 37      tbnz  w8,#0x0,LAB_1000031ac  
0000031a4 01 00 00 14      b     LAB_1000031a8
```

important section identified, it is time to analyze the assembly code. One thing to keep in mind with the disassembled code is that all the values will be stored in hexadecimal. A good place to start is with the values we know will need to be stored. In the array there are 10 variables which all get set to 0, except for the eighth array member which is changed to equal 1. In our local_dope struct there are three variables equal to 'A', '0xffffffff', and 32000. The highlighted line in the figure shows 'A' being stored in the register w8. Shortly after, #0xffffffff shows up without needing to be translated into hexadecimal. #0x7d000 is the hex for 32000. This surface level analysis offers a glimpse into the methodology of more complex reverse engineering. Lastly, the 'Hello_World!' gives us the final piece of information we need to know about what the program is.

2.3.2 CRACK ME

The next step in practicing reverse engineering is trying a Crack Me problem. Crack Me problems are toy problems designed by other software developers to help reverser's practice their skills. Usually there's an executable that opens up a little puzzle. The puzzles usually involve finding a password but can involve any number of things like decryption, key generation, etc. This section will cover the process of reverse engineering a simple crack me.

The first step I took in reversing this program was checking the header and libraries as shown in Figure 2.8 and 2.9.

Figure 2.8: Crack Me Header

```
[nataliepargas@Natalies-MBP downloads % otool -hv crackme0x00
crackme0x00:
Mach header
    magic  cpusubtype  caps      filetype ncmds sizeofcmds      flags
MH_MAGIC_64      X86_64        ALL  0x00      EXECUTE    16      1368  NOUNDEFS DY
LDLINK TWOLEVEL PIE
```

```
[nataliepargas@Natalies-MBP downloads % otool -L crackme0x00
crackme0x00:
    /usr/lib/libSystem.B.dylib (compatibility version 1.0.0, current version
1319.0.0)
```

Figure 2.9: Crack Me Libraries

Both of these steps were done with an executable dumping tool called 'otool'. The header information shows that the assembly language being used is x86. The library dump shows that the only library being linked is the system library. While these tools didn't offer that much information about this particular program, they are very useful in programs that have more dependencies and more specific assembly languages. These dumping tools ended up being vital later in the project because they offered insight into what libraries were being linked from the project's declarative language as well as what assembly language was being used with Apple's modified version of ARM64. Next, otool

was used once again to list the strings. This is a vital part of reverse engineering and is useful to be able to refer back to during the process. Figure 2.10 shows the result of the string dump.

Figure 2.10: Crack Me Strings

```
[nataliepargas@Natalies-MBP downloads % strings crackme0x00
NoxIsTheBest
Crackme Level 0x00 (created by Nox)
Enter the passphrase:
%99s
Congrats on cracking the program!!
Hmmmm maybe try again.
```

Examining this more closely, you can see that there's only a few strings that give a very clear idea of what the program is designed to do. There's a random string, the title of the program, a passphrase prompt, a string indicating the expected input, a congratulations message, and an incorrect guess message. With just this information, you could easily make an educated guess on what the password is. Most likely the passphrase will be the one random string that doesn't get given easily in the executable. But for the sake of gaining understanding of the process, it's good to go over the decompiled code. Figure 2.11 shows the decompiled code.

Figure 2.11: Crack Me Ghidra Disassembly

```

_main
entry
100003df0 55      PUSH    RBP
100003df1 48 89 e5  MOV     RBP,RSP
100003df4 48 81 ec 90 00 SUB    RSP,0x90
00 00
100003dfb 48 8b 05 fe 01 MOV    RAX,qword ptr [__stack_chk_guard]
100003dfe 00 00 00 00
100003e02 48 8b 05 ff 00 MOV    RAX=__stack_chk_guard,qword ptr [RAX]
100003e05 48 89 45 f8  MOV    qword ptr [RBP + local_10],RAX
100003e09 c7 85 7c ff ff dword ptr [RBP + local_8c],0x0
ff 00 00 00 00 00
100003e13 48 8b 05 0c 01 MOV    RAX,qword ptr [_NoxIsTheBest_100003f26]
00 00
100003e18 48 89 45 eb  MOV    qword ptr [RBP + local_1d],RAX
100003e1e 48 89 05 0a 01 00 MOV    EAX,dword ptr [_Best_100003f26+8]
00
100003e24 48 89 45 f3  MOV    dword ptr [RBP + local_15],EAX
100003e27 48 8d 05 05 01 00 AL,byte ptr [_100003f26+12]
00
100003e2d 48 85 f7  MOV    byte ptr [RBP + local_11],AL
100003e30 48 8d 3d fc 00 LEA    RDI,[s_Crackme_Level_0x00_(created_by_N_100003f33)]
00
100003e37 b0 00  MOV    AL,0x0
100003e39 e8 8c 00 00 00 CALL   _printf
100003e3e 48 8d 3d 13 01 LEA    RDI,[s_Enter_the_passphrase_100003f58]
00 00
100003e45 b0 00  MOV    AL,0x0
100003e47 e8 7e 00 00 00 CALL   _scanf
100003e4c 48 8d 7d 80 LEA    RDI=local_88,[RBP + -0x80]
100003e62 48 8d 75 eb LEA    RSI=local_1d,[RBP + -0x15]
100003e66 e8 71 00 00 00 CALL   _strcmp
100003e6b 48 85 78 ff ff MOV    dword ptr [RBP + local_90],EAX
00
100003e71 83 bd 78 ff ff CMP   dword ptr [RBP + local_90],0x0
ff 00
100003e78 0f 85 11 00 00 JNZ   Incorrect
00
100003e7e 48 8d 3d 10 00 LEA    RDI,[s_Congrats_on_cracking_the_program_100003f75]
00 00
100003e85 e8 46 00 00 00 CALL   _puts
100003e8a e9 0c 00 00 00 JMP   LAB_100003e9b

```

The first thing to do is find the entrance into the program. Ghidra already labeled the ‘`_main`’ at the top but the strings and function calls reaffirm this as the right section to be looking over. Next, thinking back to the string dump, it’s likely that the password will be near whatever function takes input. Prior knowledge of programming informs the reverser that after getting the user input, the input string will need to be compared to the right answer. Thankfully Ghidra was able to identify the function ‘`scanf`’.

Figure 2.12: Crack Me Code Section 1

```

100003e59 e8 78 00 00 00 CALL   _scanf
100003e5e 48 8d 7d 80 LEA    RDI=local_88,[RBP + -0x80]
100003e62 48 8d 75 eb LEA    RSI=local_1d,[RBP + -0x15]
100003e66 e8 71 00 00 00 CALL   _strcmp
100003e6b 48 85 78 ff ff MOV    dword ptr [RBP + local_90],EAX
ff
100003e71 83 bd 78 ff ff CMP   dword ptr [RBP + local_90],0x0
ff 00
100003e78 0f 85 11 00 00 JNZ   Incorrect
00

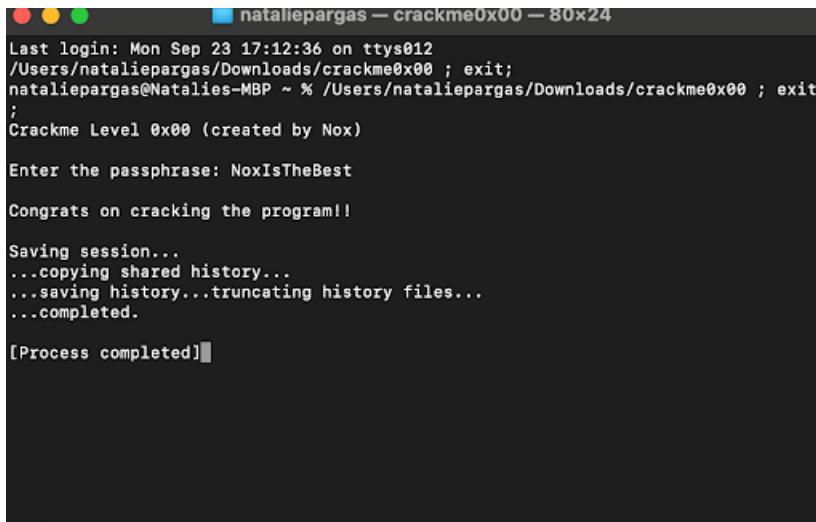
```

Inspecting the code section at Figure 2.12, it shows that the variable `local_88` contains the user input. This is because it loads a `%99s` into the `RD1` register to prepare for the input directly above where `scanf` is called. `RSI` has `local_1d` loaded into it and that variable is immediately ‘`strcmp`’ compared to `local_88`. Using these clues, one could infer that `local_1d` contains the password. To verify, a reverser should always retrace and check.

Figure 2.13: Crack Me Code Section 2

```
100003e13 TT 00 00 00 00      MOV        RAX,qword ptr [s_NoIsTheBest_100003f26] = "NoIsTheBest"
00 00
100003e1a 48 8b 0c 01      MOV        qword ptr [RBP + local_1d],RAX
```

The first instance of local_1d contains the register RAX 9. RAX has the qword ptr 's_NoIsTheBest_100003f26' moved into it right above. s_NoIsTheBest_100003f25 only contains the string "NoIsTheBest". That string is probably the password, but all that's left is to check. Figure 2.14 shows the result.

Figure 2.14: Crack Me Code Executable

```
nataliepargas -- crackme0x00 -- 80x24
Last login: Mon Sep 23 17:12:36 on ttys012
/Users/nataliepargas/Downloads/crackme0x00 ; exit;
nataliepargas@Natalies-MBP ~ % /Users/nataliepargas/Downloads/crackme0x00 ; exit
;
Crackme Level 0x00 (created by Nox)

Enter the passphrase: NoIsTheBest

Congrats on cracking the program!!

Saving session...
...copying shared history...
...saving history...truncating history files...
...completed.

[Process completed]
```

The attempt was a success and the crack me was successfully reverse engineered!

CHAPTER 3

BLUETOOTH

3.1 BLUETOOTH

Since this project involves reverse engineering a proprietary Bluetooth device, a foundational understanding of how Bluetooth operates is critical for decoding communication patterns and designing compatible interactions. Bluetooth is a widely adopted wireless technology that has seen rapid growth over the past decade. Originally developed to eliminate the need for short-range wired connections, Bluetooth is now used in countless applications—from music and video streaming to medical devices requiring continuous data updates. While it operates similarly to WiFi in some respects, the key difference is its scope: WiFi connects multiple devices to the internet, whereas Bluetooth focuses on direct device-to-device communication and offers more diverse use cases [11]. Phones, headphones, cars, watches, keyboards, and countless other devices can all make use of Bluetooth. Any hardware that needs to interface with other devices can be designed to do so via Bluetooth.

Another distinction is how Bluetooth is optimized for low-latency applications that transmit small bursts of data quickly and efficiently. This makes it ideal for scenarios where power conservation and responsiveness are more important than high throughput.

⁰Source: <https://www.klika.us/integrating-multiple-bluetooth-devices-into-a-network/>



Figure 3.1: Bluetooth Devices

3.1.1 BLUETOOTH TYPES

There are two main types of Bluetooth technology: *Bluetooth Basic Rate (BR)*, also referred to as *Bluetooth Classic*, and *Bluetooth Low Energy (BLE)*. Of the two, BLE has gained far more traction in modern applications.

Bluetooth Basic Rate is the older of the two technologies, supported in versions 1.0 through 3.0. It is primarily used for wireless audio streaming in devices like headphones, speakers, and in-car entertainment systems, and it relies on point-to-point communication.

Bluetooth Low Energy, introduced in version 4.0 and further enhanced in version 5.0, offers a broader range of capabilities. In addition to supporting audio and data transfer, BLE enables device networking and location services. It supports multiple communication topologies including point-to-point, mesh, and broadcast, making it far more versatile than its predecessor. One of BLE's standout features is its utility in high-accuracy location services. By leveraging nearby devices, BLE can determine relative positioning and provide precise location tracking.

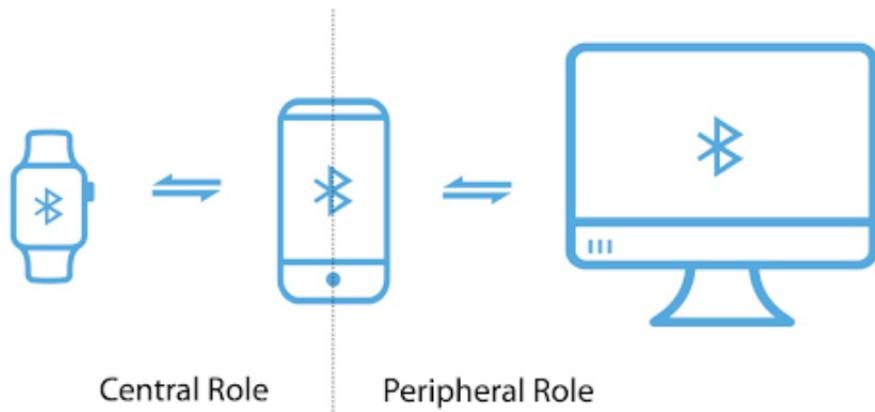
3.2 PAIRING-BASED NETWORK

The core of Bluetooth is creating a network where two (or slightly more) devices pair with each other, since Bluetooth was designed to replace corded connections. The devices being paired can be broken down into the *Central* and *Peripheral(s)*. These terms are specific to newer Bluetooth standards and can be used interchangeably with the older terms *Master* and *Slave*. In this section, I will use the newer terms.

The Central device coordinates the data being sent between it and any Peripherals. This includes handling time synchronization, sleep scheduling, and configuring the channels used through frequency hopping [14]. Peripheral devices can communicate bi-directionally with the Central device, sending or receiving data depending on the request.

Figure 3.2 shows an example of Central and Peripheral devices. Classic Bluetooth includes two network topology configurations: **Piconet** and **Scatternet**.

Figure 3.2: Central and Peripheral device relationship



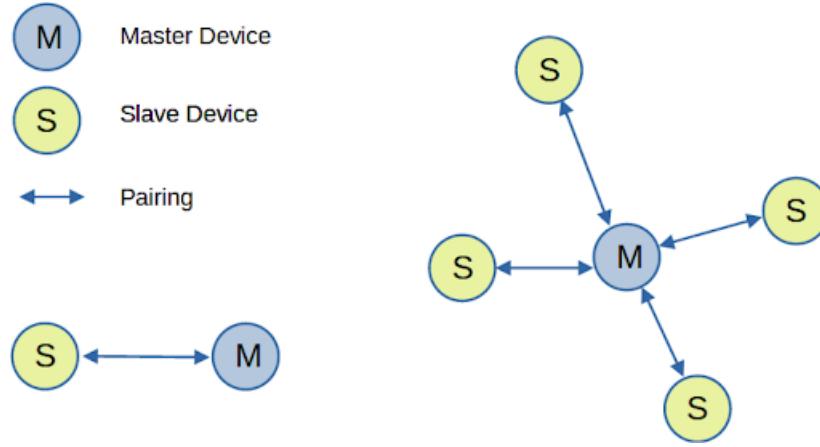
3.2.1 PICONET

A piconet is a type of ad hoc network topology that operates without preestablished infrastructure. Devices in a piconet communicate directly with one another without the need for a centralized device like a router or access point. The network is organized with

one central device and one or more peripheral devices. A single central can support up to seven active peripherals, while each peripheral can only connect to one central [14]. The central manages the timing and synchronization of all devices in the piconet. Peripheral devices cannot communicate directly with each other; the central acts as a relay, forwarding data between them.

Figure 3.3 shows the structure of a piconet.

Figure 3.3: Piconet Topology



All connections made using a BR/EDR controller occur within a piconet. BR/EDR devices communicate over the same physical channel by synchronizing with a shared clock and hopping sequence [2]. The piconet clock follows the central device's clock configuration, and the hopping sequence is also derived from the central's clock and Bluetooth device address.

Multiple independent piconets can exist in the same area. Each operates on a separate channel and is organized around a different central device. Bluetooth devices can also participate in multiple piconets simultaneously. While a Bluetooth device can only act as the central in one piconet at a time, it can serve as a peripheral in several.

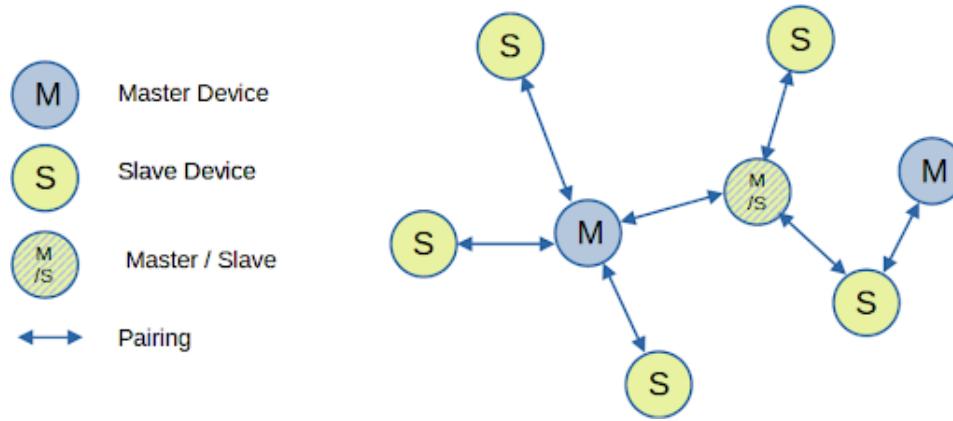
3.2.2 SCATTERNET

A scatternet is formed by connecting multiple piconets together. Unlike a piconet, a peripheral in a scatternet can connect to multiple central devices. Central devices can also connect to each other, sometimes serving dual roles as both central and peripheral.

For example, a scatternet might include a central laptop connected to a Bluetooth speaker, while also acting as a peripheral to a smartphone and printer. Figure 3.4 shows a diagram of a typical scatternet. Despite the increased complexity, peripherals still do not communicate directly with each other.

Piconets and scatternets are associated with Classic Bluetooth but are still supported in BLE and later Bluetooth versions [14]. Although BLE—with its advertising mode—has seen growing popularity, traditional Bluetooth methods still offer advantages such as added security during pairing. BLE services enhance the piconet and scatternet model, providing better security, more efficient pairing, and broader support for services.

Figure 3.4: Scatteret Topology



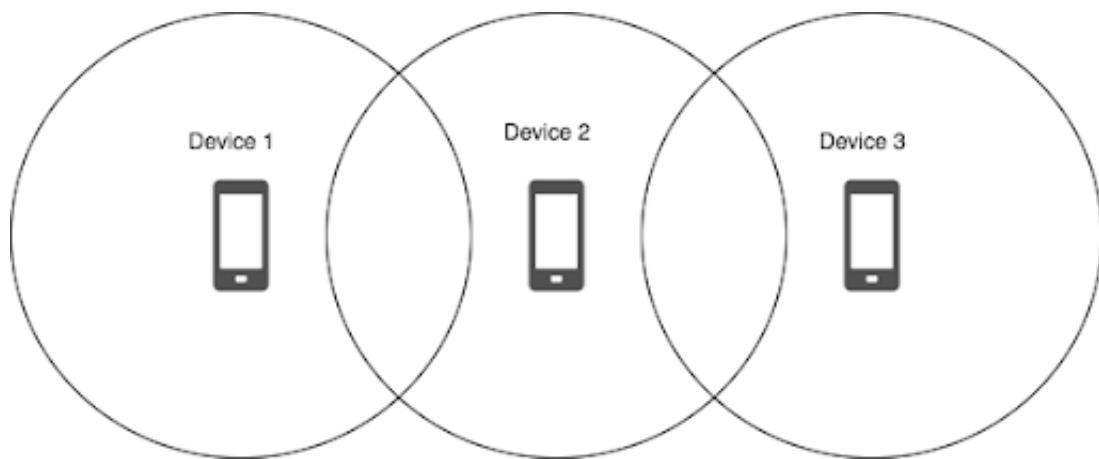
3.3 BLE MESH NETWORKS

Mesh networking, introduced with BLE, marks a significant evolution in Bluetooth topologies. In a mesh network, nodes are arranged in a self-configuring mesh, where each node relays traffic between peers when the destination is not directly connected [1].

This structure ensures a reliable path for data packets. If one path is blocked or unavailable, the network reroutes traffic dynamically. The self-organizing nature of mesh networks means they adapt in real-time to changes, such as nodes joining, leaving, or moving within the network.

For example, Figure 3.5 shows Device 1 out of range of Device 3. Device 2, within range of both, acts as a router forwarding messages between them.

Figure 3.5: Mesh topology



3.4 BROADCAST NETWORK FROM LE

BLE introduces a new feature called *advertising mode*, which operates differently from the traditional pairing between a central and peripheral. In advertising mode, a device broadcasts data openly to any BLE devices within range that are tuned to the same channel. These listening devices can receive the data and use it as needed—without any pairing, connection, or formal association.

This topology loosely resembles a multi-peripheral piconet, but with a key difference: advertising mode removes the need for a dedicated central or established links. There's no synchronization or connection—just open broadcasting.

Advertising mode has been a major factor in expanding the possibilities of Bluetooth. By

offloading the overhead of pairing, BLE made room for a wider variety of lightweight, low-latency applications. This performance shift is part of why so many modern devices—from smart beacons to virtual reality headsets—are feasible today. Devices like VR headsets, which require fast, efficient data sharing without constant handshaking, would be nearly impossible to build the same way without Bluetooth.

3.5 CONNECTIONS AND MIXED TOPOLOGY

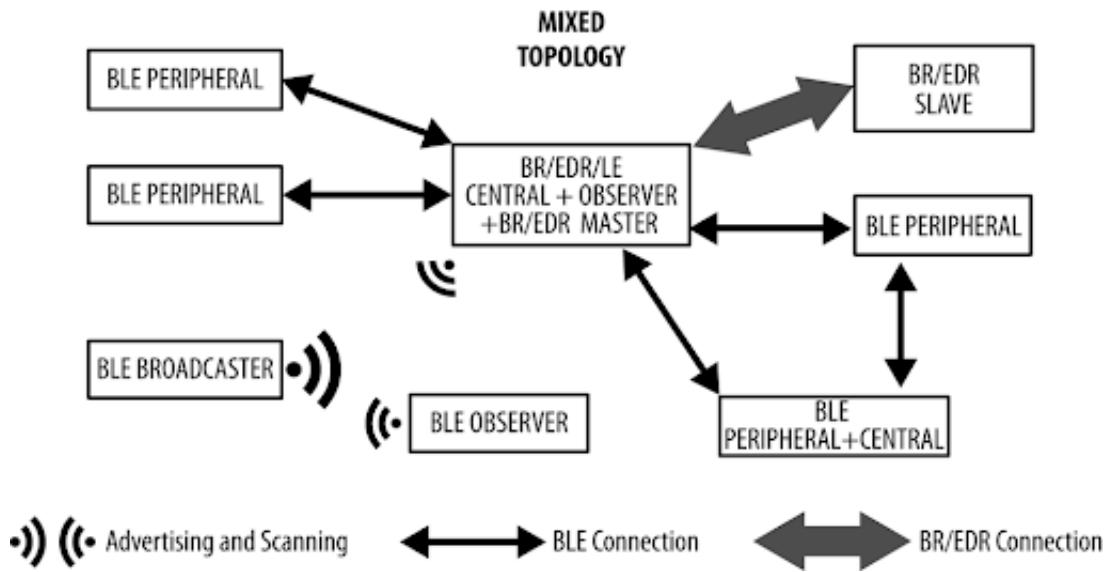
A connection becomes necessary when data need to flow in both directions or when there's more data than can fit into the two available advertising payloads. Connections in BLE are more persistent—they allow devices to remember each other and avoid reestablishing the connection every time. Once established, connections support periodic data exchanges between two devices.

This connection-based pairing is what underpins piconets and scatternets. These connections are private by design: only the two participating devices exchange data (barring external sniffing). As covered earlier, connections involve one central device and one or more peripheral devices.

BLE networks can also incorporate mixed topology configurations depending on the use case. A device that supports both BR/EDR and LE can serve as a bridge between traditional Bluetooth and BLE communication, enabling broader interoperability across protocols. The possibilities for combining these configurations are only constrained by the capabilities of the individual devices' radios and protocol stacks.

Figure 3.6 provides an example of a mixed topology Bluetooth network.

Figure 3.6: Mixed Topology



3.6 SCO AND ACL LOGICAL TRANSPORTS

With an understanding of how Bluetooth device networks are organized, we can now narrow our focus to how data is actually communicated within these networks. This communication process is crucial to keep in mind when reverse engineering a Bluetooth device.

3.6.1 BR/EDR ASYNCHRONOUS CONNECTION-ORIENTED (ACL)

ACL links are the primary way Bluetooth devices using BR/EDR exchange general data. These links handle both control signals (like LMP and L2CAP) and typical user data. ACL is *asynchronous*, meaning data is sent when available rather than on a fixed schedule.

Every Peripheral in a Bluetooth piconet gets one default ACL link to the Central. This link is established when the device first connects and is identified by a Logical Transport Address (LT_ADDR) assigned by the Central. This LT_ADDR is also used when identifying the physical connection between devices.

However, because multiple types of logical transports (like SCO links, which are used for voice) might use the same LT_ADDR, it's not enough to identify the ACL connection by LT_ADDR alone. Devices also rely on the packet type to tell them which kind of transport is being used.

ACL links can also carry isochronous data—data that needs to arrive on time, like audio streams—by setting them to automatically drop old packets. At the same time, asynchronous traffic (like file transfers) can still be sent if it's marked not to auto-flush. This means both types of traffic can share the same ACL link if configured correctly.

If the ACL link is removed or if the device loses sync with the piconet, all other connections between the Central and that Peripheral (like SCO) are also dropped immediately [2].

3.6.2 BR/EDR SYNCHRONOUS CONNECTION-ORIENTED (SCO)

SCO links are designed specifically for real-time data like voice. They create a fixed, circuit-like connection between a Central and a Peripheral by reserving slots on the Bluetooth physical channel. These connections carry 64 kbps of data, usually for audio, and are synchronized with the piconet's clock.

There are several SCO configurations that balance audio quality, latency, and bandwidth usage. Every SCO connection uses the same LT_ADDR as the ACL link, so to identify an SCO packet, a device must look at the slot number, packet type, and LT_ADDR together.

Even though SCO reserves bandwidth, the system can override these slots if higher-priority data (like control messages over ACL) needs to get through. This helps maintain overall system reliability and meet Quality of Service (QoS) requirements.

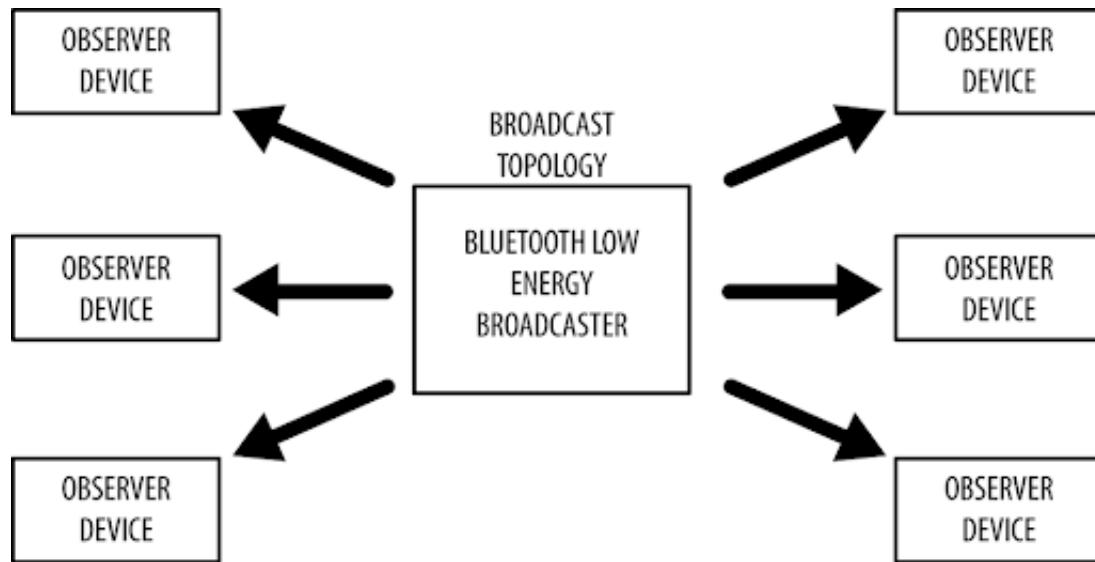
Unlike ACL links, SCO links don't carry complex protocols or layered data structures. They just send a steady stream of audio—either in a standard format or as raw data for the application to decode [2].

3.7 BLUETOOTH LOW ENERGY BROADCASTING AND OBSERVING

Bluetooth Low Energy devices communicate with the outside world in two primary ways: *broadcasting* and *connecting*. Both methods are governed by the Generic Access Profile (GAP), which outlines how devices discover and interact with each other [15].

Broadcasting enables BLE devices to send data one-way to any scanner or receiver within range. It's a form of open communication where any nearby device capable of listening can receive the broadcasted data. Figure 3.7 illustrates this communication model. In the diagram, two roles are shown: the *Broadcaster*, which periodically sends non-connectable advertising packets, and the *Observer*, which continually scans for those packets [15].

Figure 3.7: Broadcasting and Observing



Broadcasting is a key concept in BLE because it's the only mechanism that allows a device to communicate with multiple other devices simultaneously. This functionality is made possible by BLE's advertising feature.

A standard advertising packet contains a 31-byte payload, which includes information about the broadcasting device and what it offers. This data can be customized to suit different applications, providing useful context to observing devices. If 31 bytes isn't

enough, BLE allows for a secondary packet called a *Scan Response*, which provides an additional 31 bytes of payload. Observers can request this secondary data if they need more information [15].

Broadcasting is quick, efficient, and ideal for sending small amounts of data on a regular schedule to multiple receivers. However, it does come with limitations—particularly around security. Since any device within range can intercept a broadcast, there's virtually no privacy. Because of this, broadcasting isn't well-suited for transmitting sensitive or private data.

3.8 BLUETOOTH PROTOCOLS AND PROFILES

Bluetooth technology is organized into two key structural layers: protocols and profiles. Protocols form the lower-level foundation used by all Bluetooth-enabled devices. These protocols handle the encoding, decoding, packet formatting, multiplexing, and routing required for reliable wireless communication [15]. Profiles, on the other hand, are higher-level constructs that define how these underlying protocols should be used in practice. Often described as “vertical slices” of functionality, profiles ensure interoperability between devices by prescribing behavior for specific operations or applications. Profiles are typically divided into two categories: generic profiles, such as the GAP and Generic Attribute Profile (GATT), which are essential to all Bluetooth Low Energy (BLE) communications, and use-case-specific profiles, like the Proximity or Glucose Profile, which layer additional functionality on top of the generic ones to meet specialized needs.

3.8.1 GENERIC PROFILES

Generic profiles are defined by the Bluetooth Core Specification and provide the foundational mechanisms required for BLE communication. These profiles ensure that devices from different manufacturers can interoperate seamlessly. Two such profiles, GAP and

GATT, are mandatory for all BLE-compliant devices and serve as the cornerstones of BLE communication.

3.8.1.1 GENERIC ACCESS PROFILE (GAP)

The Generic Access Profile defines the roles, procedures, and operational modes required for device discovery, broadcasting, connection establishment, connection management, and security negotiation. GAP operates as the top-level control layer within the BLE protocol stack, orchestrating the basic behaviors necessary for device interaction. All BLE devices must implement and conform to the GAP specifications to ensure compatibility and reliable communication.

3.8.1.2 GENERIC ATTRIBUTE PROFILE (GATT)

GATT focuses primarily on how data is exchanged. GATT defines a universal data model and a set of procedures that enables devices to discover, read, write, and notify data values. It functions as the uppermost data layer in BLE and serves as the foundational framework upon which most BLE applications and services are constructed.

Due to their foundational roles, GAP and GATT are commonly exposed through application programming interfaces (APIs), making them the primary entry points for developers interacting with the BLE protocol stack.

3.8.2 USE-CASE-SPECIFIC PROFILES

Beyond the generic profiles, the Bluetooth Special Interest Group (SIG) has defined a series of use-case-specific profiles. These profiles are designed to standardize behavior for specific applications and are exclusively built upon the GATT framework. At the time of writing, no BLE profiles exist outside the GATT structure. However, the introduction of connection-oriented L2CAP channels in Bluetooth version 4.1 may pave the way for future GATT-less profiles.

3.8.2.1 SIG-DEFINED GATT-BASED PROFILES

The Bluetooth SIG provides an extensive catalog of standardized profiles, each tailored to support a specific application or device type. These profiles fully specify the required procedures and data formats, facilitating rapid development and ensuring interoperability across implementations.

There are several GATT-based profiles defined by the Bluetooth SIG that standardize common BLE use cases. The Find Me Profile helps locate nearby BLE devices, such as a lost phone or keyring. The Proximity Profile detects when a connected device moves beyond a certain range and can trigger alerts. The HID over GATT Profile is designed for peripherals like keyboards, mice, and remote controls, enabling the transmission of Human Interface Device (HID) data over BLE. The Glucose Profile supports the secure transmission of glucose measurement data for health monitoring. Similarly, the Health Thermometer Profile allows body temperature readings to be sent from wearable or medical devices. Lastly, the Cycling Speed and Cadence Profile enables cycling sensors to communicate speed and cadence information to compatible applications or devices.

A comprehensive list of SIG-adopted profiles is available on the Bluetooth SIG Specification Adopted Documents webpage. Developers may also consult the Bluetooth Developer Portal for up-to-date listings of adopted services and characteristics.

3.8.2.2 VENDOR-SPECIFIC PROFILES

While the SIG provides a broad range of standardized profiles, the BLE specification also permits the definition of vendor-specific profiles. These profiles are often developed to support proprietary use cases not covered by SIG standards. Vendor-specific profiles can be implemented privately between two devices (such as with a health accessory and a smartphone) or published to enable broader adoption.

Notable examples of vendor-defined profiles include profiles like Apple iBeacon which is a proprietary proximity-sensing profile used for indoor positioning and location-based

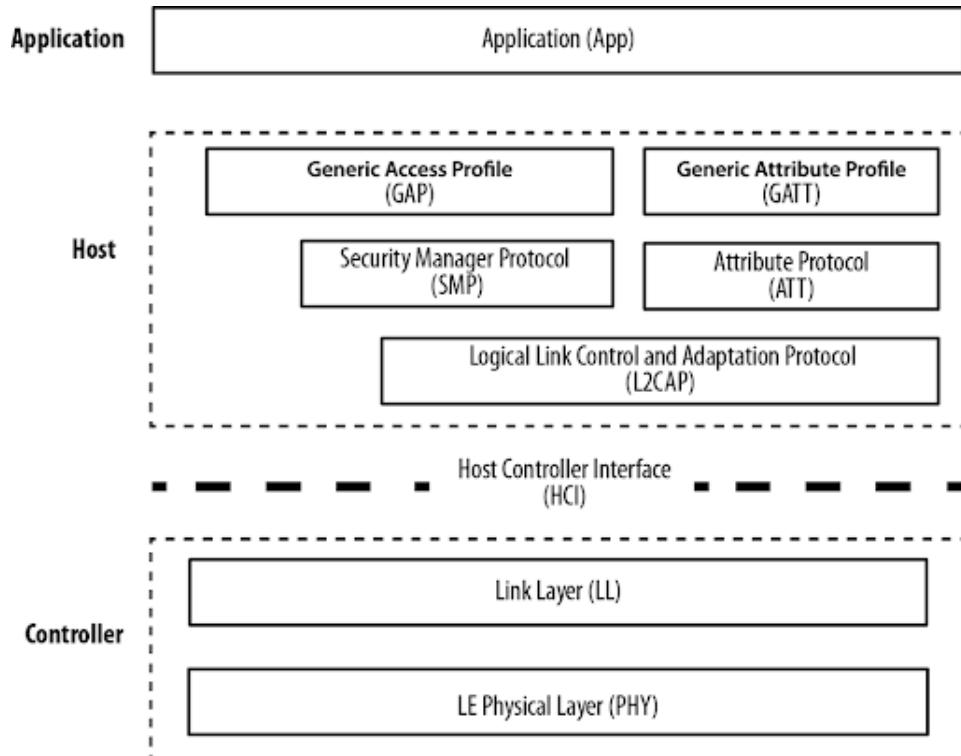
services. Another example is the Apple Notification Center Service (ANCS) which enables wearable devices or external displays to access and display iOS notifications.

These profiles demonstrate the flexibility of the BLE specification in accommodating both standardized and proprietary communication models.

3.9 BLUETOOTH LOW ENERGY (BLE) ARCHITECTURE

Although most users will only ever interact with the top layers of the Bluetooth Low Energy (BLE) protocol stack, having a basic understanding of the full stack can provide helpful context. This overview lays the groundwork for understanding how BLE devices work and why each part of the system is important. A typical single-mode BLE device is made up of three main parts: the controller, the host, and the application. Each of these parts contains multiple layers, each responsible for specific tasks that allow the device to communicate and perform its intended functions. Figure 3.8 shows the BLE protocol stack.

Figure 3.8: BLE Protocol Stack



3.9.1 APPLICATION

The application layer sits at the top of the BLE protocol stack. It includes all the logic, data handling, and user interface components that define how a device behaves in a specific use case. This layer is highly customized and varies from one implementation to another [14].

3.9.2 Host

Below the application is the host, which includes several critical layers that manage communication and device behavior.

The Bluetooth Low Energy protocol stack includes several key components, each with a specific role. The **Generic Access Profile (GAP)** controls how devices advertise themselves, discover others, establish connections, and manage roles and security [14]. The **Generic Attribute Profile (GATT)** is responsible for organizing and managing how data is exchanged between connected devices. The **Logical Link Control and Adaptation Protocol (L2CAP)** handles multiplexing and segmentation of data packets. The **Attribute Protocol (ATT)** enables data operations such as reads and writes using attribute handles. The **Security Manager (SM)** oversees pairing, authentication, and encryption processes. Lastly, the **Host Controller Interface (HCI) – Host Side** facilitates communication between the host and the controller [14].

3.9.3 CONTROLLER

At the bottom of the Bluetooth stack, the controller is responsible for handling radio operations and low-level data transport. It typically consists of several core components. The **Physical Layer (PHY)** manages the actual transmission and reception of radio signals. The **Link Layer (LL)** controls tasks such as advertising, scanning, and maintaining active connections. The **HCI (Controller Side)** works in conjunction with the host side of the Host Controller Interface to enable structured data flow between the different layers [14].

According to the Bluetooth Core Specification, a Bluetooth implementation includes a single controller, which can be configured in one of several ways. A **BR/EDR controller** contains the Radio, Baseband, Link Manager, and optionally the Host Controller Interface (HCI). An **LE controller** includes the LE PHY, Link Layer, and may also include the HCI. Alternatively, a **dual-mode controller** combines both BR/EDR and LE controller functionalities into a single integrated unit.

Together, these layers ensure that BLE devices can connect reliably and communicate efficiently. This layered architecture is often described from the bottom up — starting at the antenna and PHY layer and moving up through the stack to the user-facing application.

3.10 BLE LINK LAYER AND DEVICE ADDRESSING

In the Bluetooth Low Energy (BLE) protocol stack, the Link Layer serves as a crucial interface between the Physical Layer—which handles the actual radio transmission—and the upper protocol layers. It acts as the controller for radio operations and is essential for managing the states and timing mechanisms that define BLE communication. One of the key responsibilities of the Link Layer is to abstract the radio hardware, allowing higher-level layers to interact with it through the Host Controller Interface (HCI) [5].

In addition to timing and control, the Link Layer manages several hardware-accelerated tasks. These include generating cyclic redundancy checks (CRC) for error detection, producing random numbers for cryptographic operations, and performing encryption to secure data transfer [5]. These operations ensure the performance, integrity, and security of BLE communications.

3.10.1 LINK LAYER STATES

BLE devices operate in a well-defined set of five main states, each with specific behaviors and transitions:

- **Standby:** This is the default state, where the radio is idle and not transmitting or receiving any packets.
- **Advertising:** In this state, the device broadcasts data packets at regular intervals, making itself discoverable to other devices.
- **Scanning:** Devices in this state listen for advertising packets from others, searching for devices they might want to connect with.
- **Initiating:** When a scanning device decides to connect to an advertising device, it enters this state and sends a connection request.
- **Connected:** A persistent communication link is established. Devices in this state can exchange data regularly. The device that initiated the connection becomes the master, while the other becomes the slave [5].

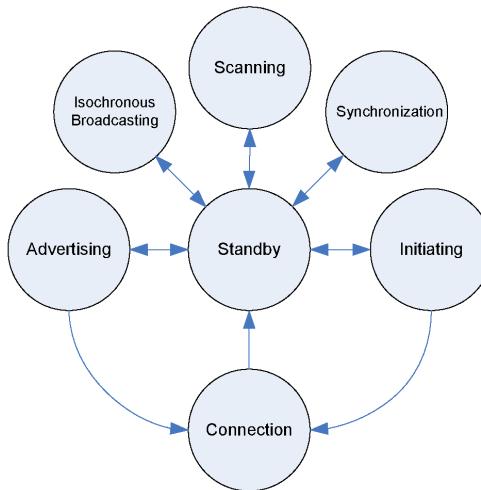


Figure 3.9: Link Layer State Machine

This framework enables BLE to support dynamic and responsive device discovery and communication, while keeping energy consumption minimal.

3.10.2 ADVERTISING, SCANNING, AND CONNECTING

Advertising and scanning are complementary operations: advertising devices periodically transmit small packets of data, while scanning devices listen for and interpret these packets. If the advertiser allows connections, and a scanning device decides to initiate one, both devices transition into the connected state [5]. These processes form the backbone of BLE interactions and are covered in more depth in subsequent chapters of most BLE technical literature.

3.10.3 BLUETOOTH DEVICE ADDRESS

Every Bluetooth device is identified by a 48-bit device address, which serves a function similar to a MAC address in networking. BLE supports two major types of device addresses: public addresses and random addresses, which influence how discoverable and trackable a device is. A public address is a globally unique, fixed identifier that is factory-programmed and registered with the IEEE.

Manufacturer:	Microsoft
Description:	Bluetooth Device (Personal Area Network)
Driver version:	10.0.26100.4061
Physical address (MAC):	44:FA:66:B5:A3:4E

Figure 3.10: Example Bluetooth Address

Because of this registration requirement, it ensures uniqueness across all BLE devices worldwide. However, it can also present privacy concerns, since it does not change and is easily traceable [5]. Random addresses, by contrast, provide more flexibility and are often preferred in consumer devices to protect user privacy. These addresses can either be static or private, and they do not require IEEE registration.

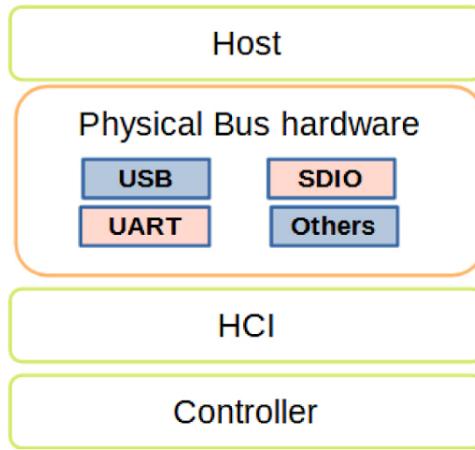
Static Addresses are used as a drop-in replacement for public addresses. A static address remains the same between sessions but can be regenerated at boot or kept for the lifetime of the device until a power cycle occurs.

Private Addresses are designed to change regularly to enhance privacy. There are two subtypes of private addresses, **Non-resolvable private addresses** and **Resolvable private addresses**. Non-resolvable private addresses are temporary, randomly generated, and cannot be traced back to a specific device. They are rarely used due to limited functionality. Resolvable private addresses use an Identity Resolving Key (IRK) and a random number to generate temporary addresses. These change over time, making the device hard to track by unknown parties. However, trusted or bonded devices that have the IRK stored can resolve these addresses, allowing them to re-identify the device securely [5].

This system allows BLE to balance two key goals: persistent connectivity with trusted devices and protection against unwanted tracking or eavesdropping by unknown entities.

3.11 HOST CONTROLLER INTERFACE (HCI) LAYER

The Host Controller Interface (HCI) serves as the communication bridge between the host and controller layers in a Bluetooth system. It is what enables communications between the lower layers of the bluetooth stack. These components can be implemented on the same chipset or on separate chipsets. When they are on separate chipsets, HCI is critical for enabling interoperability and communication between them. In such cases, HCI defines both the standardized protocol and the physical transport mechanisms that facilitate message exchange. As shown in Figure 3.11, the host, HCI, and controller are depicted as distinct components. HCI commonly uses transport technologies such as UART, USB, and SDIO to establish physical connections between the host and controller. When the host and controller are integrated on the same chipset, HCI functions as a logical interface rather than a physical one [5].

Figure 3.11: Host, HCI, and Controller components

3.11.1 UNIVERSAL ASYNCHRONOUS Rx/Tx (UART)

UART is a serial communication protocol that uses wired connections to exchange data between devices. It is commonly found in integrated circuits and is especially favored in devices like mobile phones and laptops, where the entire Bluetooth stack is implemented on a single chipset. UART is also well-suited for small, resource-constrained devices powered by microcontrollers, as it is lightweight and efficient in its use of system resources [14].

In the context of the HCI UART protocol, four types of packet transmissions are defined: the Command packet (identified by 0x01), Event packet (0x04), Asynchronous Connection-Less (ACL) packet (0x02), and Synchronous Connection-Oriented (SCO) packet (0x03). Because the receiving HCI layer cannot distinguish these packet types based on content alone, each transmission begins with a leading byte — known as an indicator — which acts as a header to identify the type of packet that follows [14].

3.12 CONNECTION

To establish a connection in Bluetooth Low Energy (BLE), the central device begins by scanning for advertising peripheral devices that are currently accepting connection requests. Advertising packets can be filtered based on the Bluetooth Address or the contents of the advertising data itself. Once a suitable advertiser is detected, the master sends a connection request packet. If the slave responds, a connection is established [15].

The connection request packet includes several critical parameters that define how communication will proceed. One such parameter is the frequency hop increment, which determines the hopping sequence used throughout the connection. In addition, the packet specifies three connection-related parameters. The connection interval defines the time between the start of two consecutive connection events, ranging from as short as 7.5 milliseconds for high throughput to as long as 4 seconds for reduced power consumption. The slave latency indicates how many connection events the slave device is permitted to skip without risking disconnection. Lastly, the connection supervision timeout sets the maximum duration allowed between the receipt of valid packets before the connection is considered lost [15].

Once connected, communication happens through repeated connection events, which are periodic time slots where the master and slave take turns exchanging packets. A connection event always begins with a packet from the central, requires the peripheral to respond if it receives a packet, continues until both sides have no more data to send, is repeated at the connection interval until the connection is either closed or lost, and can be closed by either the master or the slave. If the central sends a packet and does not receive a response, it waits until the next connection event to resume. To manage interference and ensure privacy or security, BLE supports a white list mechanism at the Link Layer. This list defines which Bluetooth device addresses are of interest. Devices not on the list are ignored — their advertising (by scanners) or connection request packets (by advertisers) are simply dropped [15].

3.13 SERVICES AND CHARACTERISTICS

In Bluetooth Low Energy, services and characteristics define how data is structured, communicated, and managed on a server. A clear understanding of this architecture is essential for tracing data flows and identifying commands in any BLE-enabled device.

3.13.1 ATTRIBUTE PROTOCOL (ATT)

The Attribute Protocol (ATT) defines how a server exposes its data to a client, and how that data is organized and accessed in a Bluetooth Low Energy (BLE) system [5].

3.13.1.1 ROLES IN ATT

There are two main roles in ATT communication:

- **Server:** The server is the device that stores and exposes data, and may allow certain behaviors to be remotely controlled. It receives commands from peer devices and sends responses, notifications, or indications in return. For example, a thermometer acts as a server when it provides access to its current temperature, unit of measurement, battery level, or the interval at which it records readings. Instead of requiring the client to repeatedly poll for changes, the server can proactively notify the client when data updates occur [5].
- **Client:** The client is the device that reads data from the server or controls the server's behavior. It sends commands and requests, and accepts incoming notifications and indications. In the thermometer example, a mobile device that connects to the thermometer and reads temperature values is operating as the client [5].

3.13.1.2 DATA FORMAT: ATTRIBUTES

The data exposed by the server is organized into attributes. An attribute is a general term for any piece of data available on the server, such as services or characteristics (which are described further in the text).

Each attribute in the Bluetooth Attribute Protocol (ATT) is composed of several key elements. The first is the attribute type, identified by a Universally Unique Identifier (UUID), which distinguishes the kind of data the attribute represents. A 16-bit UUID is used for standard attributes adopted by the Bluetooth SIG, while a 128-bit UUID is reserved for custom or vendor-specific attributes defined by developers [5]. The second element is the attribute handle, a 16-bit value assigned by the server that functions like an address. This handle uniquely identifies each attribute during the connection and allows the client to reference it directly. Handle values range from 0x0001 to 0xFFFF, with 0x0000 being reserved [5]. Lastly, attribute permissions determine what operations—such as read, write, notify, or indicate—are allowed on the attribute and what security conditions must be met. These permissions are defined outside of the ATT layer, typically in the GATT or in the application logic [5].

3.14 PROFILES

3.14.1 GENERIC ATTRIBUTE PROFILE (GATT)

The GATT defines how BLE devices organize, expose, and exchange data over a connection. It builds directly on top of the Attribute Protocol (ATT), using it as the transport layer for transmitting structured data known as attributes [5].

GATT introduces a hierarchical data model centered around three key concepts: Services, Characteristics, and Profiles [5].

These elements allow BLE devices to encapsulate user-facing data and device functionality in a standardized, interoperable way. All BLE application-level interactions happen within the framework defined by GATT [15].

3.14.2 GATT ROLES

Like ATT, GATT supports client and server roles, which are dynamic per transaction rather than fixed per device. This means a single device can simultaneously act as a GATT server for one connection and a GATT client in another [5].

- **GATT Server:** Stores and exposes attributes to the client. It responds to requests and may send unsolicited updates via notifications or indications. Every BLE device must at least support a minimal GATT server, even if only to respond with error codes [15].
- **GATT Client:** Initiates communication by performing service discovery, reading/writing attributes, and subscribing to updates. The client has no prior knowledge of the server's structure—it learns through discovery procedures [15].

3.14.3 SERVICES

A service is a logical grouping of one or more attributes that collectively provide a certain function on the server. A service always includes at least one characteristic, and often contains supporting attributes like declarations, descriptors, and included services (used to reference other services) [5].

There are two main classifications for services. There are **Primary Services** which represent core functionality for a Bluetooth device. Services such as Battery Service or Device Information Service are primary services. Primary services are more commonly used than the other type of service, **Secondary Services**. Secondary services offer auxiliary support and must be referenced by primary services [5].

The Bluetooth SIG maintains a set of adopted services with published specifications to ensure interoperability across vendors. If a product claims compliance with one of these services, it must strictly follow its specification [5].

3.14.4 CHARACTERISTICS

A characteristic is the smallest logical unit of user-accessible data on a BLE server. Characteristics always belong to a service. They also include values which contain the actual data being exposed, properties which are the operations permitted to be performed on the value (e.g., read, write, notify, indicate), and metadata about the value (e.g., format, units, configuration settings).

For example, the battery level characteristic within the Battery Service lets clients read the device's current power level [5].

3.14.5 PROFILES

While services and characteristics define how data is stored and exposed on the server, profiles describe how clients and servers interact, including service usage, connection procedures, and security requirements. Profiles are not discovered over BLE connections like services are; instead, they exist as specification documents, often adopted by the Bluetooth SIG. These define how multiple services should be used together to achieve specific use cases (e.g., Heart Rate Profile, Blood Pressure Profile) [5].

Each profile specification includes: required services and characteristics, interaction behaviors for both server and client, connection and security requirements, and example usage diagrams and workflows [5].

3.15 NORDIC SEMICONDUCTOR

3.15.1 nRF CONNECT

nRF Connect for Mobile—formerly known as the nRF Master Control Panel—is a feature-rich application designed for developers working with Bluetooth Low Energy (BLE) devices. Created by Nordic Semiconductor, this free app is available on both Android and iOS platforms and is widely regarded as one of the most powerful tools for BLE development and testing [1].

The application supports numerous Bluetooth SIG-adopted profiles, including Over-the-Air Device Firmware Updates (OTA DFU), making it especially useful for firmware deployment and debugging [1]. Key features include BLE device scanning with real-time RSSI graph visualization, advertisement data parsing for both general and manufacturer-specific payloads, and comprehensive GATT client functionality that allows connection to peripherals and inspection of their services and characteristics. Users can interact with characteristics by performing read, write, and observe operations, and can enable notifications and indications to receive real-time updates from connected devices. The application also supports GATT server emulation with built-in presets, OTA firmware updates for supported devices, and automated testing through XML-defined test scripts to validate BLE device behavior. Additionally, it provides UART service support for communication over UART-over-BLE channels [1].

It is compatible with Android version 4.3 and higher or iOS version 8.0 and higher. At the time of writing, the Android version supports more functionality than the iPhone version.

With its comprehensive feature set, nRF Connect is an essential tool for BLE developers, from prototyping and debugging to final deployment [1].

CHAPTER 4

ANDROID REVERSE ENGINEERING

Java is the basis of Android applications and has historically been what is used. Java is particularly well-suited for reverse engineering due to its architecture. It is designed to be platform-independent by running on the Java Virtual Machine (JVM), which functions as an abstract computing environment with its own instruction set and memory model. The JVM interprets a binary format known as a class file, which contains bytecode (JVM instructions), a symbol table, and metadata.

The JVM is stack-oriented, with most operations involving the operand stack in the current execution frame. New frames are created whenever methods are invoked, providing a clear and modular structure. This design simplifies the decompilation and analysis process, which is critical when reconstructing application logic and behavior [13]. Figure 4.1 illustrates the flow of data from Java source code through compilation, class loading, bytecode verification, and execution by the JVM.

The use of class files is one of the key factors that makes reverse engineering Java-based applications more straightforward. Unlike many other languages that are compiled directly into low-level assembly, Java is compiled into bytecode, an intermediate representation that retains a significant amount of the original structure of the source code. This structure provides helpful context for decompilation tools. As noted by Alex Kalinovsky, "bytecode can be interpreted or compiled after loading, which results in a two-step transformation of the high-level programming language into the low-level machine code. It is the intermediate step that makes decompiling Java bytecode nearly flawless" [12]. Bytecode

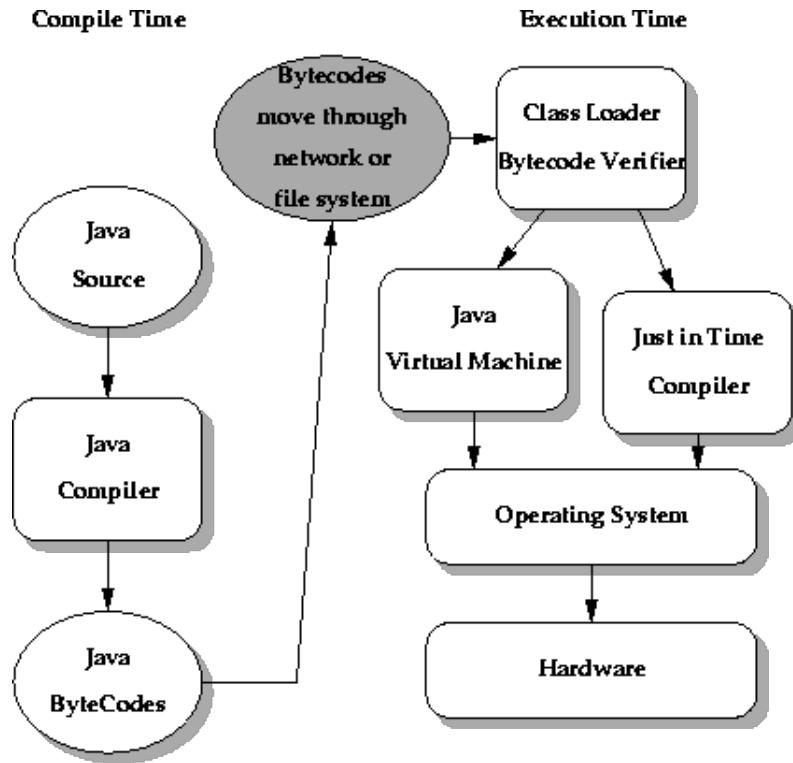


Figure 4.1: Java Source to Hardware

preserves nearly all critical information from the source file, with the exception of comments and formatting. Method names, variables, and control flow are typically intact, making it easier to reconstruct the original logic. Additionally, the stack-oriented design the JVM offers a predictable execution model that further aids analysis.

A man-in-the-middle (MITM) approach utilizing instrumentation was selected for its relative simplicity compared to full decompilation or bytecode-level modification. In this context, the MITM or instrumentation method refers specifically to intercepting communication between the client application and the device it controls (e.g., a host device such as a Bluetooth-connected peripheral). This strategy includes techniques such as using Bluetooth sniffers to capture transmitted data and injecting code at runtime to observe or alter behavior without directly modifying the underlying APK or its bytecode.

4.1 A BRIEF OVERVIEW OF ANDROID FUNDAMENTALS

To understand the reverse engineering process of an Android application, it's helpful to start with some foundational knowledge. At the most basic level, an Android app is made up of activities and layouts. Activities are special classes, typically written in Kotlin, that control how the app behaves and responds to user input. For example, if an app has a button, the activity defines what happens when that button is clicked.

Most Android apps have one or more screens, and the design of these screens is handled through layout files or directly in the activity code. Layout files are usually written in XML and can include elements like text, images, and buttons [8]. Figure 4.2 shows a diagram of activities and layouts. Figure 4.3 illustrates how they interact during actual device use.

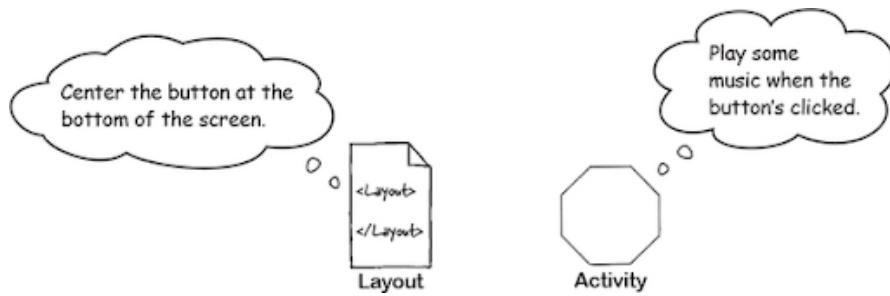


Figure 4.2: Android Activity Layout

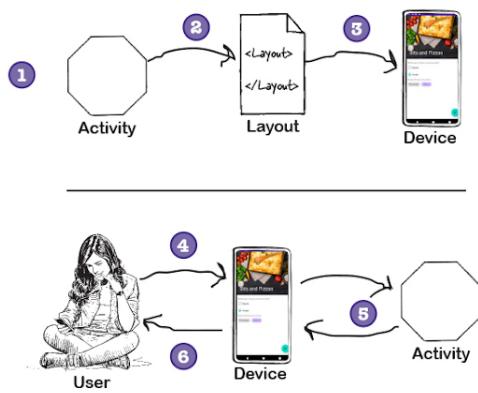


Figure 4.3: Android Activity to User

- Android launches the app's main activity.

- The activity instructs Android to use a specific layout.
- The layout is displayed on the device.
- The user interacts with the layout.
- The activity responds to these interactions and updates the display.
- The user is able to view the updated display seamlessly.

Apps are built using the Android Software Development Kit. The most straightforward place to do this is the Android Studio app. The Android SDK has Android source files and a compiler used to compile code into the Android format.

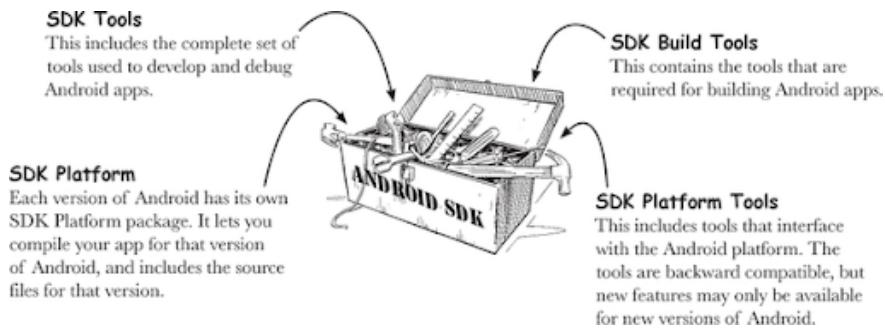


Figure 4.4: Androiod SDK Tools

If you're familiar with Android, you've probably encountered its fun, food-based versioning system. Each Android version includes a version number and a code name. The version number refers to the specific release (like 8.0), while the code name represents a broader release range (like Oreo). Google stopped using dessert-themed code names after Android 9. Each version of Android also corresponds to an API level, which is used by apps to ensure compatibility. For example, Android 11 corresponds to API level 30. Android Studio projects use the Gradle build system to compile and deploy apps. Gradle-based projects follow a standard directory structure, which is shown in Figure 4.5.

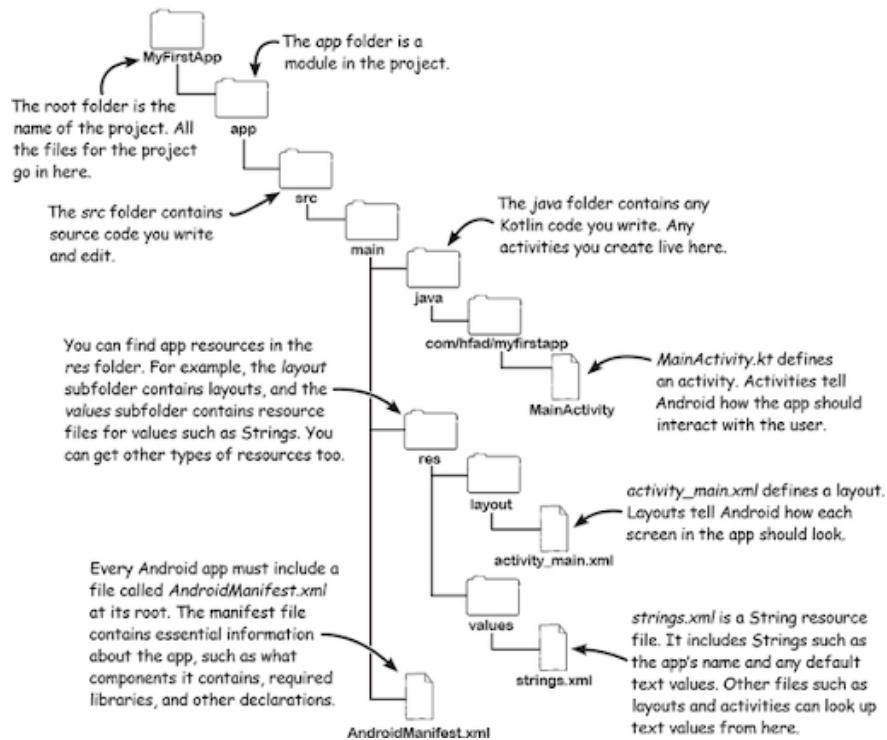


Figure 4.5: Gradle Build Structure

A much larger version of this structure is seen in the decompilation of the Persolips app.

An Android app is made up of four main components: activities, services, broadcast receivers, and content providers. Each of these acts as an entry point into the app, either for the user or the system. Some components operate independently, while others rely on or interact with each other. As mentioned before, Activities are the entry points that users interact with directly. Each activity represents a single screen with a user interface. In a note-taking app, for example, you might have one activity for composing a note, another for creating folders, and another for reading existing notes. While each activity functions independently, they work together to provide a cohesive user experience. Activities can even be started by other apps—if permissions allow. For instance, the Photos app might open a note composition activity so an image can be saved directly into a note.

Activities also manage key interactions between the system and the app. They help Android determine what's important to the user and ensure that relevant processes are

kept running. They track backgrounded processes that might be reopened and help manage killed processes so their state can be restored if needed. Activities enable apps to work together and let the system coordinate those interactions. In code, activities are implemented as subclasses of Android's Activity class.

Services are also entry points, but they run in the background and are designed for general-purpose tasks rather than direct user interaction. They're used to keep apps running behind the scenes—for example, to perform long-running operations or interact with remote processes. A service has no user interface. It might do something like keep the flashlight on while a user switches to another app or fetch data from the internet without interrupting the user's experience. Activities and other components can start services and allow them to run independently or bind to them.

There are two types of services: started services and bound services. A started service continues to run even if the user leaves the app—like keeping the flashlight on or syncing data in the background. Services that the user is aware of (like the flashlight) are handled differently than those that operate silently (like data syncing). Bound services, on the other hand, only run when another component is actively connected to them.

Broadcast receivers allow the system to deliver messages or events to apps, even if the app isn't currently running. This makes it possible to respond to system-wide events outside of the usual user flow. For example, a scheduled notification could be delivered through a broadcast receiver without needing the app to stay open. Content providers manage access to data stored outside of the app itself, such as in a local SQLite database, on the device, or in the cloud. They make it possible for apps to query or modify this data securely. For example, a social networking app might use a content provider to access the user's device contacts when adding new friends.

When an Android app is packaged and deployed, it's bundled into a file called an APK (Android Package Kit). This file contains everything needed to install and run the app on a device. Figure 4.6 outlines the typical structure of an APK. Kotlin source code is compiled

into bytecode, and libraries and resources are gathered. These elements are then assembled into the APK. The contents of an APK typically include:

4.1.1 STRUCTURE OF AN APK FILE

An APK file is a ZIP archive typically containing:

- **META-INF/**: Contains cryptographic metadata.
 - **MANIFEST.MF**: Lists files with SHA-1 digests.
 - **CERT.SF**, certificate files: Signatures ensuring integrity.
- **lib/**: Native libraries compiled per architecture (e.g., `armeabi-v7a`, `arm64-v8a`, `x86`, `x86_64`).
- **res/**: Application resources (images, layouts), not compiled into **resources.arsc**.
- **assets/**: Raw asset files accessible via **AssetManager**.
- **AndroidManifest.xml**: App metadata in binary XML (can be decoded using tools like Apktool or Androguard).
- **classes.dex**: Bytecode in Dalvik Executable (DEX) format for the Android Runtime.
- **resources.arsc**: Precompiled binary resources (e.g., strings, styles).

All of these components work together to form the complete, installable app package. When installed, the Android system unpacks the APK and sets up everything the app needs to run on the device.

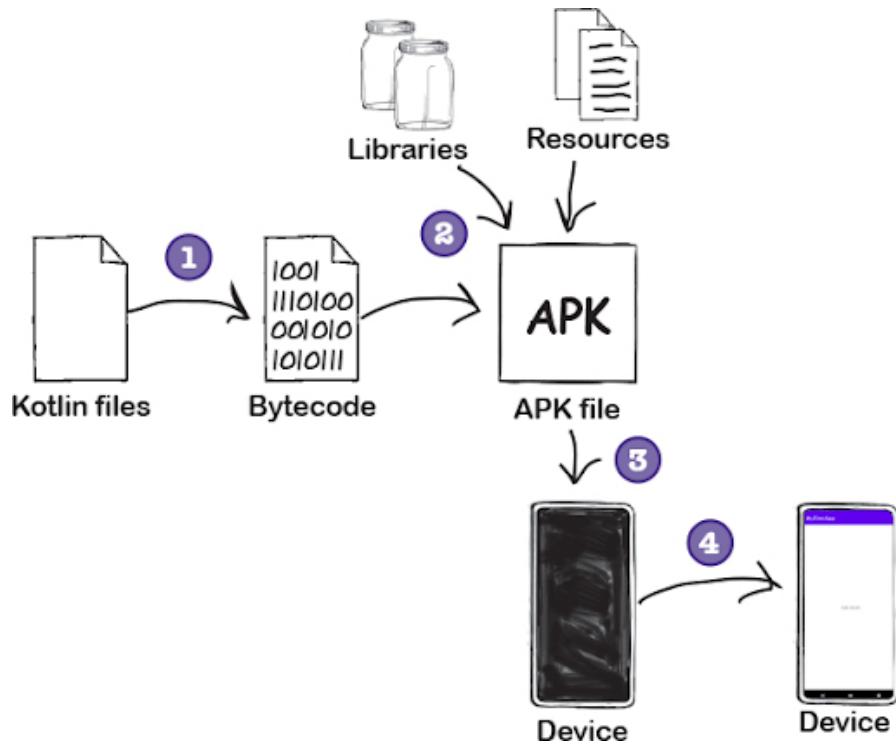


Figure 4.6: Android App Components

4.2 OVERVIEW OF THE YVES SAINT LAURENT MAKEUP PRINTER (ROUGE SUR MESURE CUSTOM LIP COLOR CREATOR)

One of the most basic yet crucial steps in reverse engineering a device is acquiring a comprehensive understanding of its functionality. The YSL makeup printer is a unique device, and it is therefore essential to examine its components in order to contextualize any observed or modified behavior.

At its core, the device functions as a lipstick dispenser capable of mixing precise quantities of pigment to produce custom colors. It operates in tandem with the Rouge Sur Mesure application. Within the app, users can select colors in one of three ways: shade matching from photographs, choosing from a color wheel, or employing YSL's shade stylist, which uses outfit analysis to recommend a complementary lip color. When executing a dispense cycle, the interface progresses through three primary screens. Initially,

it displays the selected shade, followed by a progress screen featuring three dots and the label of the inserted color family, and finally the device presents the successfully printed color. According to YSL, “Rouge Sur Mesure uses custom color cartridges and artificial intelligence to make and recommend over 7,000 unique lip colors”¹. The system leverages PERSO technology, L’Oréal’s latest innovation in personalized beauty, to analyze user data and tailor product recommendations to each individual. Figure 4.7 presents an image of the device, Figure 4.8 illustrates the color wheel selection interface, and Figure 4.9 shows the image-based shade matching screen.



Figure 4.7: Rouge Sur Mesure Custom Lip Color Creator

¹https://www.yslbeautyus.com/rouge-sur-mesure.html?srsltid=AfmBOorRGtCnqu9Hg_-9mebHE1hqicByITz-JtpICKBy8BZVciVP-ml



Figure 4.8: Color Wheel App Screen



Figure 4.9: Color from Image App Screen

Despite the promise of limitless customization, the device's actual capabilities are more constrained. It contains three slots into which cartridges from one of seven color families: red, nude, orange, pink, warm red, warm nude, and cool nude may be inserted. Figure 4.10 displays the color wheel grouped by family, and Figure 4.11 depicts the physical cartridges, including the cool nude sample used in this analysis. Each color family must be purchased separately, and at the time of writing, the cost per family is approximately \$65 USD.



Figure 4.10: Rouge Sur Mesure Color Families



Figure 4.11: Rouge Sur Mesure Cartridge Trios

These limitations reveal that the system does not fully eliminate the need for multiple lipstick purchases. Although the app can generate custom shades within a single color family, the presence of three distinct "nude" families underscores that users must still invest in multiple cartridges to achieve a broad spectrum of tones. Additionally, many users

report that the printer struggles to reproduce an exact match for the selected custom color. In my own testing, the device seldom produces an exact duplicate of the chosen shade; instead, it defaults to the closest available hue within the loaded family. To encourage additional purchases, the application will notify users if a slightly closer match exists in an unloaded color family, prompting them to acquire that family.

Understanding the functionality of the device overall is essential for correlating the data captured during reverse engineering with its operational components.

4.3 REVERSE ENGINEERING THE ROUGE SUR MESURE APP

Makeup enthusiasts are all too familiar with the struggle of purchasing numerous lipsticks that differ only slightly in shade, all in pursuit of achieving a specific look. One common scenario involves finding inspiration online, seeing a color that looks stunning on a model, and feeling compelled to replicate it, only for the shade to appear entirely different in person. Another frequent challenge is the search for the perfect "your lips but better" shade, which often leads to buying fifteen different options that are eventually discarded because the undertones are too peachy or the shades lean too cool. Both situations reflect a pattern of consumption that is costly and wasteful, yet there remains no clear way to avoid it.

This is the problem that the Rouge Sur Mesure Custom Lip Color Creator by YSL attempts to address. By integrating technology with beauty, the device is designed to allow users to custom-mix lip colors and match nearly any shade. However, the product has received numerous complaints regarding its functionality. Users have specifically criticized the way lipstick colors are grouped, which limits the range of shades the device can produce.

The goal of this project is to investigate the internal workings of the device in detail, to the extent that I can demonstrate my understanding by altering how a dispense occurs without relying on the official app. By gaining a deeper understanding of how the device

operates, I hope to improve its capabilities and tailor its behavior to better suit my own needs.

4.4 ENVIRONMENT SETUP

A variety of tools were used to support both static and dynamic reverse engineering throughout the project. Each tool was selected for its reliability, popularity in the reverse engineering community, and compatibility with Android-based systems.

4.4.1 ROOTING THE PHONE

Rooting the OnePlus 8T was a necessary step for me in the broader context of reverse engineering a mobile application. I had several reasons for wanting to root the device, all of which stemmed from the technical demands of the project I'm working on.

The first and most pressing reason is that rooting is essential for hooking into the app's internal processes. Through my early experimentation, I discovered that modifying the app's behavior would require a Man-In-The-Middle (MITM) attack, and tools commonly used for this—such as Frida—typically require root access to function properly.

Additionally, rooting a device grants access to lower-level Bluetooth packet data, which is crucial for my analysis. On a rooted phone, I can use Frida scripts to log system-level Bluetooth reads and writes that would otherwise be inaccessible. If progress slows using the current toolset, more advanced offensive techniques, such as those available through Kali NetHunter or Ubertooth, may become necessary. These tools also work more effectively on a rooted device.

Rooting a phone, however, wipes all data stored on it. For this reason, any data I intend to keep must be backed up and restored after rooting is complete. Overall, reverse engineering is an inherently unpredictable process, and having access to a wide range of

tools—many of which require root access—means I can pivot more quickly if one approach hits a dead end.

The first hurdle in rooting the device was downgrading the operating system. Although I found a guide for rooting the phone on Android 14 (the version it shipped with), it required a specific boot image that I couldn't locate. My phone's variant, the OnePlus 8T KB2005 (global version), receives incremental boot image updates, which means a complete boot image for patching and reuploading isn't publicly available.

I attempted to use a boot image I found online, but it didn't match my Android version, which resulted in the phone becoming soft-bricked. To resolve this and simplify the patching process, I used the MSM Download Tool to downgrade the device to Android 11. I followed the XDA Developers guide: [XDA MSM Tool Guide](#)².

After connecting the phone to my PC via a trusted USB port, I had to troubleshoot some driver issues. Interestingly, the front USB ports on my PC didn't recognize the phone reliably, while the back ports—connected directly to the motherboard—worked better. When the device appeared as QHUSB_BULK, I knew the Qualcomm USB drivers were improperly installed. I downloaded the correct drivers from Microsoft's catalog: [Qualcomm USB Drivers](#).

Next, I confirmed my build version: kb2005_14.0.0.602(EX01), indicating I had the International model, which uses the KB05AA firmware. I downloaded the corresponding MSM tool and followed these steps: Launch MsmDownloadTool V4.0.exe.

1. At the login screen, choose **Other** and click **Next**.
2. Select the correct target region (**O2** for Global).
3. Press **Start** to initiate the flashing process.
4. Power off the device and allow it to cool.

²<https://xdaforums.com/t/op8t-oos-kb05aa-ba-da-unbrick-tool-to-restore-your-device-to-oxygenos.4180837/>

5. Enter **EDL (Emergency Download)** mode by holding both volume buttons and plugging the phone into the computer.
6. Wait for the flash to complete (approximately 300 seconds).

With the phone successfully downgraded, I proceeded to root it using another guide from XDA³. I began by re-enabling Developer Mode, which involved tapping the build number eight times in the settings. Once Developer Mode was active, I enabled USB Debugging and plugged the device back into my computer. To ensure everything was working properly, I ran **adb devices** and **fastboot devices** to confirm the device was being recognized. If drivers weren't functioning correctly, I had to reinstall them.

Next, I needed to identify which slot the device was currently using. I did this with the command **fastboot getvar all**, which showed that the current slot was a. With that information, I moved forward with extracting the boot image from the active partition. I used a limited functionality TWRP recovery image that I booted into temporarily with **fastboot boot recovery.img**. TWRP stands for Team Win Recovery Project, and it's an open-source project that provides custom recovery images for Android devices. It's designed to replace the default recovery screen that comes with the phone, giving access to features that aren't available in the stock version. These features include installing custom ROMs, creating full device backups, restoring data, wiping partitions, and more. This version of TWRP doesn't include a GUI but provides ADB shell access, which was all I needed. Inside the shell, I used the **dd** command to copy the **boot_a** partition to the SD card. After exiting the shell, I pulled the **boot_a.img** file from the device using ADB and then rebooted the phone normally.

Once I had the **boot_a.img** file on my computer, I transferred it to the phone's internal storage, placing it in an accessible location like the Downloads folder. I then installed the latest version of the Magisk Canary APK on the phone. Opening the app, I selected the

³<https://xdaforums.com/t/oneplus-8t-all-variants-root-magisk-oxygenos14-oos14.4703449/>

install option and chose "Select and Patch a File," pointing Magisk to the boot_a.img I had extracted earlier. This generated a patched image file named magisk_patched_a.img, which I copied back to my computer.

To proceed with rooting, I rebooted the phone into fastboot mode using adb reboot bootloader. Instead of flashing the patched image, I temporarily booted into it using the command fastboot boot magisk_patched_a.img. This gave me temporary root access. With the device booted into this patched environment, I opened Magisk again and selected the install option, this time choosing "Direct Install (Recommended)" to apply root access permanently to the internal boot image. After the installation completed, I rebooted the phone and used a root checker app to verify that root access was working. Everything checked out, and the device was successfully rooted.

4.4.2 VIRTUAL MACHINE

To ensure that the reverse engineering process does not compromise the host system, it is advisable to work within a virtual machine (VM). This not only adds a layer of protection against potentially harmful software, but also provides the flexibility to use tools that may not be compatible with the native operating system. In this case, setting up a virtual machine was necessary for both reasons. Since the focus was on reverse engineering the Android version of the app, many essential tools—such as those for extracting APKs or interfacing with the device—required a Windows environment. Rooting the phone, for example, involved installing legacy drivers from outdated and potentially unsafe sources, increasing the risk of introducing malware. Using a VM offered a controlled environment where such risks could be isolated.

Because the target phone needed to be connected externally via USB, a virtual machine with USB passthrough support was required. Initially, a macOS system was the only available host, so Parallels Desktop was selected. USB passthrough functionality had recently been released in version 20.3.0 (May 7, 2025) (source), just shortly before the virtual

machine was set up. However, issues arose when attempting to get the device recognized within the VM. Given that the phone was an older model that also had trouble with driver installation on native Windows machines, the problem was likely not due to the VM itself. It is also worth noting that Parallels requires a paid license, which should be considered when evaluating virtualization options.

Access to a native Windows machine later resolved many of the earlier issues. Windows-based VMs tend to perform more reliably on Windows hosts, and the improvement in stability was noticeable. VMware Workstation was chosen for this setup. It supports USB passthrough effectively and is free to use with a Broadcom account.

4.4.2.1 JADX

The primary decompiler used for this project was JADX⁴, a widely adopted tool specifically designed for reverse engineering Java-based Android applications. JADX takes an APK or XAPK file and converts the included .dex (Dalvik Executable) files—Android’s equivalent of Java .class files—into readable Java source code by generating .jar files. While the decompiled Java output is generally highly readable, it is typically not suitable for direct recompilation or execution without additional modification. Nonetheless, it serves as a valuable resource for understanding app structure, logic, and flow.

4.4.2.2 FRIDA

Frida is a dynamic instrumentation toolkit used to inject JavaScript snippets or custom libraries into native processes. Instrumentation can be defined as modifying a program’s control flow by adding probes or hooks into said program. It leverages QuickJS to run scripts inside a target process, providing full memory access, function hooking, and the ability to call native functions at runtime. Frida establishes a bi-directional communication channel between the injected script and the host environment, enabling real-time monitoring and

⁴<https://github.com/skylot/jadx>

manipulation of application behavior. In this project, Frida was primarily used to extract runtime information from the app and to override or modify methods on the fly.

4.4.2.3 WIRESHARK

Wireshark is a widely used network protocol analyzer, typically applied to internet and server traffic. In the context of this project, it was used to analyze Bluetooth communication. Although Wireshark does not natively support capturing Bluetooth packets without additional hardware, it becomes highly effective when paired with a Bluetooth sniffing device. Additionally, it can display Bluetooth logs exported from Android devices with Bluetooth HCI snoop logging enabled. This made it an essential tool for inspecting communication between the app and the connected device.

4.5 ALTERING THE APP

To ensure that I would be able to successfully complete this project, I first needed to confirm that I could alter the app's behavior in some meaningful way. At the time, the only tools I had available were my MacBook with Android Studio installed. I was running the app on an emulated Android device and quickly found out that in order to use Frida—the dynamic instrumentation toolkit I planned to rely on—I needed to root the emulator. While I had decompiled parts of the app using JADX, modifying the code that way wasn't practical. Decompiled code generally can't be recompiled cleanly, and reconstructing the original logic is both difficult and time-consuming. Given that Frida was going to be the foundation of my approach for real modifications later on, it made the most sense to start testing with it from the beginning. While it's true that .apk files for this app weren't obfuscated and could have been manually edited and re-signed, that path would have diverged from the strategy I'd ultimately use. Instead, I committed to building out a working Frida workflow right away.

The first step was gathering all necessary tools. This included Frida itself, the frida-server binary compiled for android-arm64 (version 16.7.14), and Objection, a utility built on top of Frida for runtime mobile app exploration. I also installed the Android command line tools directly from Google's site, which are separate from Android Studio's full IDE.

With those in place, I turned to setting up a rooted Android emulator. I used rootAVD, a script that modifies emulators to grant root access. Since I was working on an ARM64 Mac, I selected version 12 (Android S) from rootAVD's compatibility chart. Using Android Studio, I downloaded the system image for API level 28 with an ARM64 architecture. After confirming the image was downloaded, I launched the emulator from the terminal using the emulator command. I first listed the available virtual devices, then ran the appropriate command to launch the desired AVD.

Inside the rootAVD directory, I ran the script by typing `./rootAVD.sh`. I listed all the AVDs and selected the correct ramdisk.img from the path `system-images/android-28/default/arm64-v8a`. If everything was patched successfully, the emulator would shut down and then boot up again normally. A grey or black screen during boot was a sign that something had gone wrong. Once the emulator was running again, I verified root access using ADB. I ran `adb shell`, then typed `su`, followed by `whoami`. If the output returned "root," that meant the emulator was successfully rooted.

At that point, I was ready to run Objection. I launched it in a new terminal window by attaching it to the app's process using the command `objection -g com.loreal.ysl.perso.lips explore`. If everything was working correctly, Objection would connect to the app, allowing for runtime exploration and testing.

Next, I needed to get frida-server running on the device. I navigated to the folder where the binary was located and made it executable with `chmod +x`. Then I pushed it to the emulator's file system using ADB. Specifically, I placed it in `/data/local/tmp` using the command `adb push frida-server-16.7.14-android-arm64 /data/local/tmp/frida-server`. After

entering the device shell and switching to root again with su, I launched the server in the background using `./frida-server &`. A successful start would return a process ID.

With the Frida server running and everything hooked up, I tested a simple proof-of-concept string replacement. I began by choosing the string I wanted to replace: “No activities yet.” To work with it at the memory level, I converted the string to hexadecimal using `echo -n "No activities yet" | xxd -p`, which gave me the raw hex representation of the text. I then used Frida’s memory search capabilities to locate that hex sequence in the app’s memory.

Next, I decided on the replacement string—“Natalie activities”—and repeated the process to convert it to hex using the same `xxd` command. Then I retrieved the process ID (PID) of the running app using `adb shell pidof com.loreal.ysl.perso.lips`. With that PID, I ran the Frida script I had written to replace the original string with the new one. The command was `frida -U -p <PID> -l replace_textview.js`. Occasionally, Frida throws unhelpful syntax errors—like missing semicolons—that don’t actually affect runtime behavior. Even if such an error appeared, I would go back to the app, let it load, and check if the text had been replaced. And in this case, once the relevant screen rendered, the string had indeed been swapped out, confirming that the setup was working.

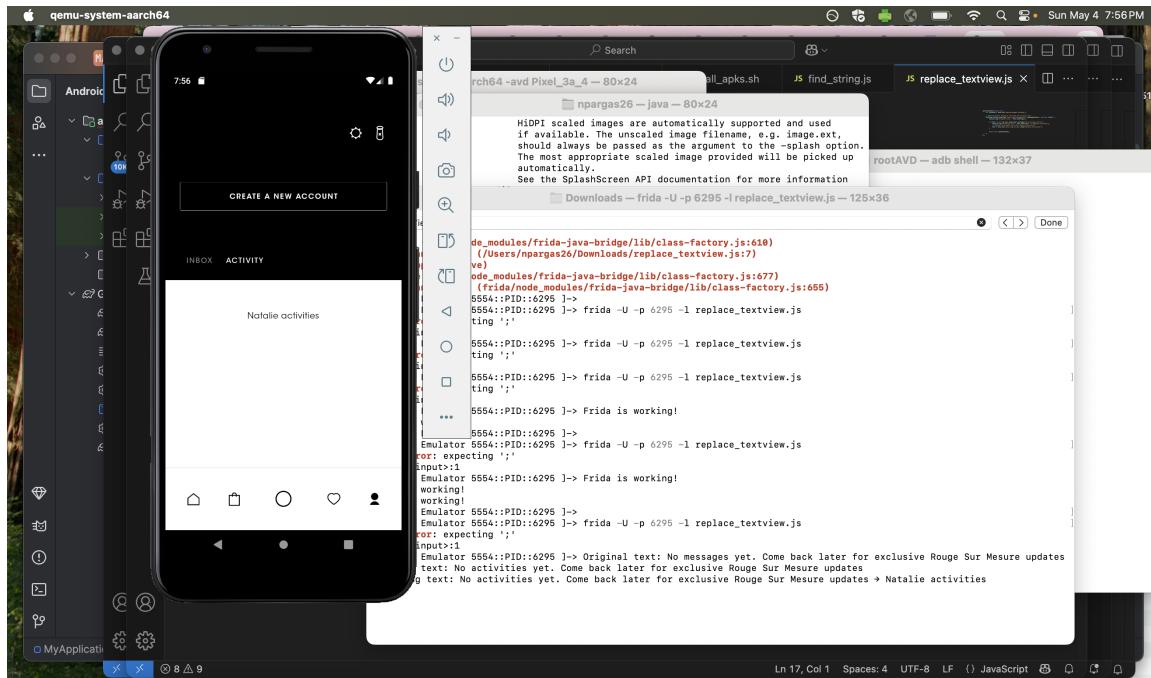


Figure 4.12: Altered App Screen

CHAPTER 5

METHODOLOGY

Reverse engineering is a complex process with no universal solution. It often involves extensive trial and error, making it important to adopt a strategy that encourages rapid iteration while minimizing risk. In this case, the most practical approach was to prioritize simplicity and accessibility. The Android version of the app was selected for analysis, as it uses Java and Kotlin—languages that are generally more amenable to reverse engineering. Additionally, a man-in-the-middle (MITM) technique was chosen to intercept and manipulate data exchanged between the app and the target device.

The initial step involved verifying the ability to interact with the app by modifying data in a detectable way. Once this was confirmed, the next objective was to capture the data being transmitted between the app and the makeup printing device, along with identifying the method of transmission. The final step was to replicate that communication in order to send custom data to the device in the expected format.

5.1 DECODING APP DATA

When trying to understand the communications between the app and the makeup printer device, one method that's certifiably free of interpretation error is capturing the packet data sent over the air between the devices. Understanding the raw packets is also helpful if the app itself is too difficult to reverse engineer, as it enables direct interaction with the device and bypasses the app altogether. This is a common reverse engineering method because

developers seldom bother to obfuscate Bluetooth packet data. Ultimately, the relevant data has to appear somewhere in the Bluetooth packets, it's just a matter of decoding where and how it does.

5.1.1 TESTING WITH THE nRF CONNECT APP

There are a few different ways to capture this packet data. As mentioned previously, one method is using a Bluetooth packet sniffer along with the nRF Connect app and Wireshark. This was my first test. The following figures show the device information I discovered using the nRF Connect app. It was fascinating to see how much information was accessible through nRF Connect alone. For example, Figure 5.3 shows the NUS_ID, NUS_RX, and NUS_TX identifiers, which I previously found in Jadx after a long period of searching. The nRF Connect app revealed this information immediately, without even needing to connect to the device.

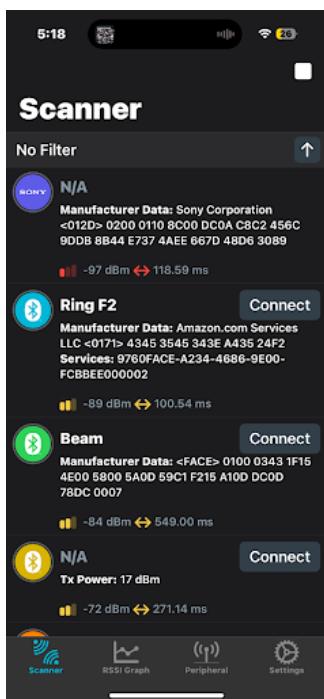


Figure 5.1: nRF Connect Scanner page

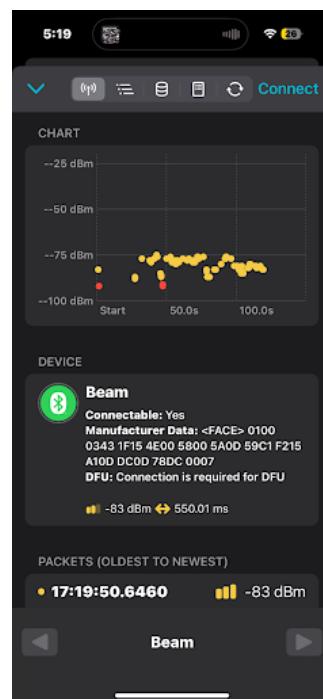


Figure 5.2: nRF Connect Beam Beacon page



Figure 5.3: nRF Connect Beam Services and Characteristics

Figure 5.4: NRF Connect screenshots

```
/* loaded from: classes2.dex */
public class BleManager extends ScanCallback implements TransportLayerDelegate {
    private static final int MTU = 200;
    private static final UUID NUS_ID = UUID.fromString("6E400001-B5A3-F393-E0A9-E50E24DCCA9E");
    private static final UUID NUS_RX = UUID.fromString("6E400002-B5A3-F393-E0A9-E50E24DCCA9E");
    private static final UUID NUS_TX = UUID.fromString("6E400003-B5A3-F393-E0A9-E50E24DCCA9E");
    private static final long SCAN_PERIOD = 5000;
    private static final String TAG = "Ble";
```

Figure 5.5: Jadx App Gatt Profile Information

Unfortunately, while this information was helpful, I wasn't able to gather much more from nRF Connect because the device would connect and then disconnect shortly afterward. Even so, this discovery was valuable because it suggests that a handshake occurs between the app and device after connecting, which the nRF Connect app does not perform.

Without dumping the firmware, I cannot know exactly what occurs during this handshake. However, I can begin piecing it together by examining the onConnectionStateChanged() function in Jadx. Figure 5.6 presents a model of this function. It illustrates how the function checks Android's BluetoothProfiles to determine whether the constant indicates a disconnected state (0) or a connected state (2). If connected, the function proceeds to request an MTU of 200 from the device. This is likely the point at which issues arise, as the app's custom MTU is significantly larger than the general Bluetooth MTU used by the nRF Connect app.

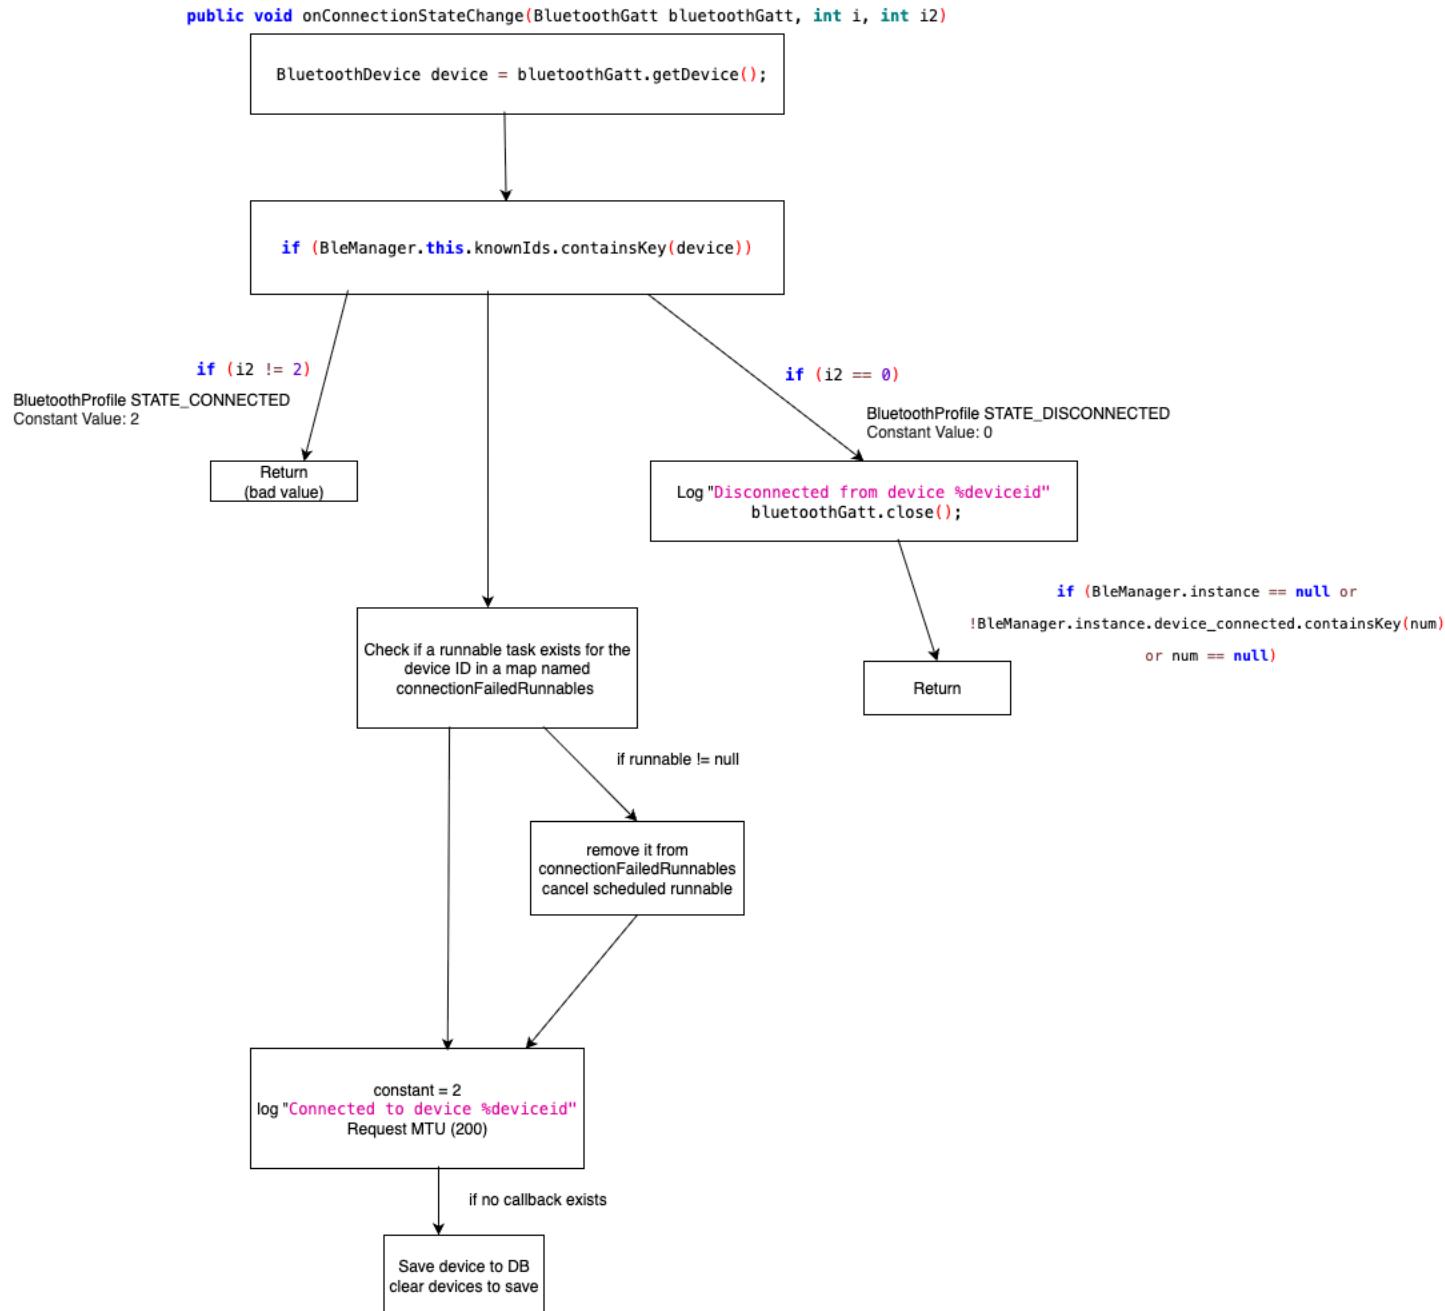


Figure 5.6: `onConnectionStateChange()` State Machine

Nevertheless, if the device had remained connected, it would have been possible to conduct tests and observe the raw byte payloads along with the characteristics transmitting them. For certain devices, it is even feasible to reverse engineer them using only the nRF Connect app, which is both impressive and somewhat concerning.

After discovering the connectivity issue on iOS with the nRF app, I was able to access the device through the nRF Connect app by switching to Android. I had the nRF Connect app running when I opened the YSL app and connected to the device, which triggered a popup on my phone indicating that an “unknown device was detected” and asking whether I wanted to connect for debugging.

When I reopened the nRF Connect app, the device and all of its packet information became visible. However, when I returned to the YSL app, the device appeared as disconnected and required me to disconnect it in the nRF Connect app before I could reestablish connectivity. Using this workaround, I was able to run several tests and capture definitive live packet data. During these tests, the nRF Connect app attempted to translate the packet data into Unicode strings. This translation revealed that the shade information is transmitted by the device in one of the initial packets, regardless of whether a dispense occurs or a color is selected.

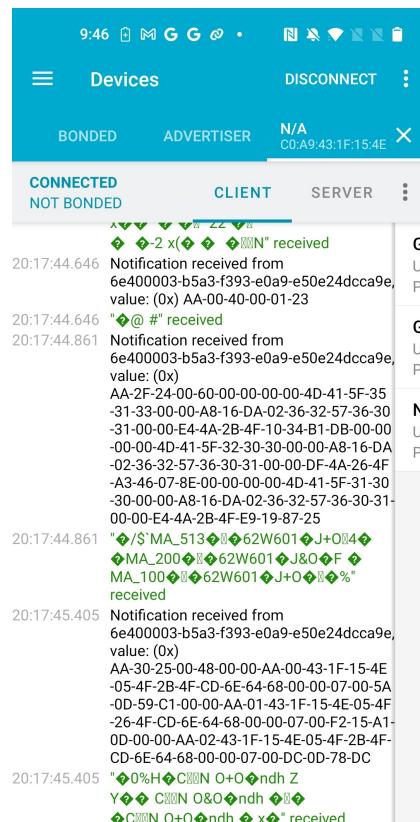


Figure 5.7: nRF Connect Android Device Communication

5.2 TESTING WITH THE NRF SNIFFER AND WIRESHARK

In addition to using the nRF Connect app, a Bluetooth sniffer proved to be especially useful. While nRF Connect can capture communications received from the target device, it doesn't show all Bluetooth traffic. Bluetooth sniffers, on the other hand, can capture the full set of Bluetooth commands exchanged between devices. The main drawback of relying solely on a sniffer is the lack of built-in packet decoding or interpretation. Using both tools together offers the most complete picture of the communication.

When selecting a sniffer, I had to weigh the overall cost of setting up the environment, which led me to choose the nRF52840 Dongle. My research indicated that this dongle was reliable and offered significant value for its modest \$10 cost. Additionally, it is manufactured by the same company that appeared to provide the software development kit used by the target device. Although more advanced and expensive sniffers can yield more detailed results, I recommend that anyone beginning this process start with a reasonably priced sniffer and consider investing in a higher-end tool only if the initial setup proves insufficient.

Setting up the nRF52840 Dongle required several steps. The dongle itself is simply a piece of hardware that serves no specific function until the appropriate software is flashed onto it. The first step involved installing the nRF Util tool, which depends on the presence of the SEGGER J-Link software. If this software is missing, outdated, or incorrectly configured in the system path, the flashing process will not work properly. The next requirement was to install the ble-sniffer command. Completing these installations prepared the environment for flashing the dongle.

The subsequent step involved locating the directory where nRF Util was installed. Within this directory, the corresponding firmware for the sniffer could be found, usually named in the format `sniffer_nrf52840dongle_nrf52840_*.zip`. After connecting the dongle to my computer—I use a MacBook equipped only with USB-C ports, though a USB adapter functioned without issue—I proceeded to install the nRF Util device command using `nrfutil`

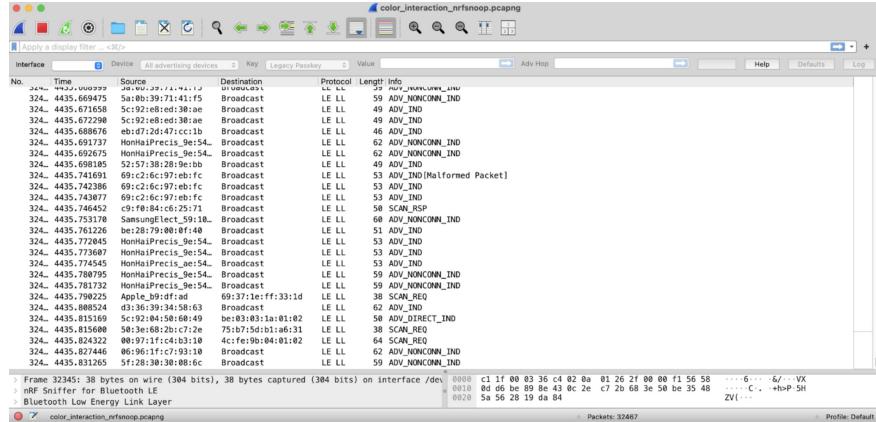


Figure 5.8: Wireshark Output from Bluetooth Packet Sniffing

install device. Running the command `nrfutil device list` allowed me to identify the correct device. Once detected, I programmed the dongle by executing `nrfutil device` program with the options **-serial-number <serial-number> -firmware <file>**, substituting the serial number identified earlier and the appropriate firmware file. At this stage, the dongle was flashed and ready for use.

Following these steps, I was able to begin sniffing Bluetooth traffic immediately. When the dongle successfully captured packets, an LED light on the device would blink with each sniffed packet. The setup functioned perfectly on the first day. However, during retesting the following day, the LED stopped blinking, and the dongle was not recognized correctly. Diagnosing the problem required extensive troubleshooting. Sometimes, the dongle appeared in the nRF Connect app, while at other times it disappeared entirely and was not visible from the command line. Ultimately, I had to wipe the dongle and repeat the flashing process, which seemed to restore its functionality.

Nevertheless, I encountered a significant issue: the dongle worked a bit too well. I was testing in a densely populated apartment complex where numerous Bluetooth devices were active, as previously observed through the nRF Connect app. The dongle's LED blinked incessantly, making it impossible to isolate packets relevant only to my target device and specific tests. Figure 5.8 illustrates a log in which the dongle captured over 32,400 packets.

Although Wireshark provides powerful commands for filtering captured data, there was insufficient information to determine suitable filters for isolating the target packets. Attempts to filter based on segments of UUIDs yielded no matches, and manually parsing such a large volume of data was impractical. Fortunately, with an Android device, it is possible to extract HCI snoop logs, which record all of the phone's Bluetooth interactions. To reduce the volume of data and improve the chances of identifying relevant packets, the next step was to pull these logs directly from the phone.

5.3 ANALYZING THE HCI SNOOP LOGS

The first step in obtaining Bluetooth HCI snoop logs is to enable Bluetooth debugging on the device. This requires ensuring that USB debugging is already activated, which should be the case given the extensive ADB debugging performed on the phone thus far. Once confirmed, navigate to the developer settings and enable the option labeled "Enable Bluetooth HCI snoop log."

With logging enabled, it is then necessary to generate traffic to capture in the log. For my initial test, I chose to record a connection event between the application and the makeup printer device. Retrieving the log from the device proved somewhat more complicated than anticipated, as the file was not stored in the expected location. To locate it, I ran the following command: **find /data -iname '*btsnoop*'**. This is a prudent step to avoid attempting to pull the log from an incorrect path. Ultimately, the btsnoop log was located in the directory: **/data/misc/bluedroid**.

Bluedroid is Android's Bluetooth stack, and files stored in its directories are inaccessible to non-root users due to the system's permissions model. To retrieve the log, I first copied it to the device's SD card using the command:

```
cp /data/misc/bluedroid/2025_05_28_14_24_03_btsnoop.txt /sdcard/btsnoop_hci.log
```

I then verified the file's presence on the SD card by executing **ls -l /sdcard/btsnoop_hci.log**.

Once confirmed, I used the adb pull command to transfer the log to my Mac for further analysis.

With the log successfully copied to the computer, it became ready for examination in Wireshark. To contextualize the data, it is useful to recall the role of Bluetooth's Host Controller Interface (HCI), which defines the standard communication protocol between the higher-level host and the lower-level controller. Generally, these logs are text files that capture HCI commands, events, and data packets. As a brief reminder, HCI commands are instructions sent from the host to the Bluetooth controller (for example, a disconnect command), while HCI events are the corresponding responses sent from the controller back to the host (for example, a disconnection complete event). Each entry in the log represents a specific communication action or response.

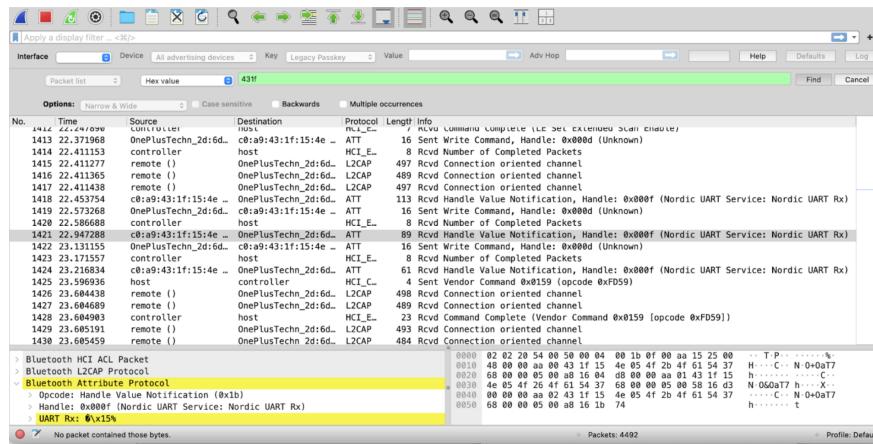


Figure 5.9: Wireshark Output from HCI Snoop Logs

Despite filtering the capture to a specific event, the log still contained a substantial number of packets (4,482 in total). To simplify analysis, it is advantageous to use filters within Wireshark. One useful filter for this scenario is the btatt filter, which allows the user to isolate only the Attribute Protocol (ATT) traffic. This protocol governs the discovery, reading, and writing of attributes between a client and a server device.

By narrowing the captured data using this filter, it becomes significantly easier to identify relevant packets that represent interactions between the phone and the device. For

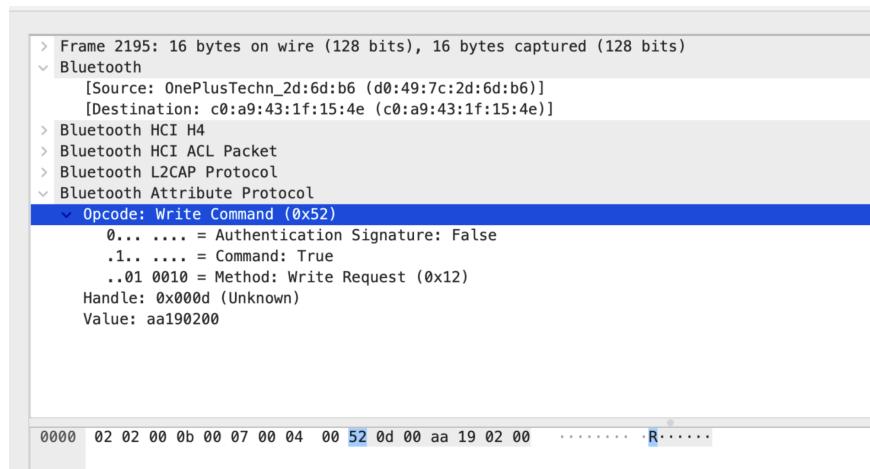


Figure 5.10: HCI Packet Tree in Wireshark

instance, in Wireshark I identified communication from the phone, OnePlusTechn_2d:6d:b6, directed to the device c0:a9:43:1f:15:4e, in which the phone issued a write command.

By examining the attribute tree in Wireshark, it is possible to locate both the handle and the value of this write operation. In this case, the write occurred to the characteristic with Handle: 0x000d, and the value written was aa190200. To gain a comprehensive understanding of each service and characteristic involved, this analysis process can be repeated for each interaction of interest within the application.

The primary advantage of this method over others is that it reveals low-level Bluetooth interactions occurring on the phone, as well as the bidirectional communication between the device and the phone. Throughout this project, I have found that the ability to dynamically employ various investigative methods as specific challenges arise is critical to the decoding process.

While the information captured in these logs is invaluable, the complexity of the device means that decoding packet data alone would be an excessively time-consuming approach for understanding or modifying the communication. To gain insight into how these packets are constructed and sent, it is necessary to examine the internal functions of the application itself. For this purpose, the next step is to utilize Frida to instrument and analyze the relevant application code.

5.4 SCRIPTING

Frida accomadates the use of scripts to perform specific functions in conjunction with their application. It is a very useful tool in the reverse engineering process to be able to target and extract specific information. The first script created was one to enumerate the contents of the BLE send function. This function is how any and all BLE communications get broadcast through the application so it holds a wealth of information in terms of figuring out how the device is expecting to receive data. While it is hypothetically possible to hook into the Android library's Bluetooth send function, it is much more difficult to do with less payoff than hooking into the native send function. The send() method was discovered by reading through the decompiled code in jadx. Its existence was verified by listing the methods in the BLEManager class using Frida.

```
Java.perform(function () {  
    var BleManager = Java.use("com.vinsol.loreal.PersoLips.utils.BleManager");  
});
```

Frida has some of its own terminology. `Java.perform()` is one of the most common Frida methods because it alerts Frida where to hook into the application. Here, Frida is being instructed to hook into the `BleManager` class.

```
var originalSend = BleManager.send.overload('[B', 'int');
```

This line saves the original `send(byte[], int)` function so it can be called as necessary. It will need to be called as usual at the end of enumeration so that there isn't an interruption in the methods.

```
function enumerateObject(obj) {  
    try {  
        var objClass = obj.getClass();  
        console.log("Object class: " + objClass.getName());
```

```
var fields = objClass.getDeclaredFields();
for (var i = 0; i < fields.length; i++) {
    var field = fields[i];
    field.setAccessible(true);
    try {
        var name = field.getName();
        var value = field.get(obj);
        console.log(" Name: " + name + " = " + value);
    } catch (err) {
        console.log(" Could not access field: " + err);
    }
}
} catch (err) {
    console.log("Error during enumeration: " + err);
}
}
```

This is the definition for the function to enumerate the contents of each object in the `send()` function. The `try{}/catch{}` allows the function to execute safely without crashing. The `getClass` Java method gets the class of the object at runtime. It then loops through each field and sets them as accessible. For each field in the loop, it logs the name and object value.

```
BleManager.send.overload('[B', 'int').implementation = function (bArr, i) {console.log
```

Here is where the `send` function actually gets overridden. It starts by logging that the `send` function was called.

```
var bytes = Java.array('byte', bArr);
```

```
var hex = Array.from(bytes).map(function (b) {  
    return ('0' + (b & 0xFF).toString(16)).slice(-2);  
}).join(' ');\nconsole.log("Payload (hex): " + hex);
```

This code segment is where the byte array output of `send` gets converted into hex to make it easier to read. It starts by turning the Java byte array `bArr` into a format that can be read by Frida. The bytes are then ANDed with `0xFF` to convert signed bytes to unsigned. It then converts the result into base 16 (hex). Finally, `slice(-2)` ensures each value is represented by two characters (zero-padded if needed).

```
console.log("Enumerating fields of BleManager instance...");  
enumerateObject(this);
```

This logs and inspects the current `BleManager` instance — the object calling `send()`. It gives insight into the app's internal state at the moment of the call.

```
return originalSend.call(this, bArr, i);
```

Finally, the original `send()` function is called to allow normal app behavior to continue uninterrupted.

5.5 ANALYZING BLE PAYLOADS VIA FRIDA HOOKING AND PACKET INSPECTION

When using Frida for packet analysis, the first and most crucial step is ensuring that I can capture the packet data along with any functions I intend to hook. Fortunately, this is achievable directly within Frida by identifying where Android's `writeCharacteristic` function is implemented. This function is responsible for transmitting data over Bluetooth Low Energy (BLE). By using JADX, I was able to locate this function without much difficulty

in the BleManager class, as discussed in the scripting section. Within this class, the relevant method is appropriately named `send()`.

With the app's `send()` method hooked, I was able to perform various tests where I made small alterations to the dispense output to observe how these changes affected the raw payload. I experimented with modifications to color, volume, and dispense time, hoping to identify clear RGB values within the packet alongside the corresponding volume information for each print command. Based on research into other color-based BLE devices, such as LED lights, these values often appear clearly in some portion of the payload. Unfortunately, this device proved to be more complex. Given the novelty of the technology and the significant investment behind its production, this complexity is not surprising. While the Bluetooth communications were not encrypted, they were intricate enough that simply knowing what the device was supposed to do was not sufficient; a deeper understanding of the app's source code was necessary.

One drawback of hooking the `send()` function instead of using a finely tuned Bluetooth sniffer was that I was only able to capture the communications that the hook managed to log. This excluded packets received from the device. Additionally, the timing of these logged transmissions might not be entirely accurate since they are not captured immediately but rather passed through the hook. Although I had access to a Bluetooth packet sniffer, the output was overly verbose, filled with external Bluetooth noise, and ultimately too time-consuming to filter within the constraints of my project. Time constraints often lead to the abandonment of reverse engineering efforts, so given the number of time-intensive tasks involved, I opted to prioritize monitoring the `send()` function even at the cost of some fidelity.

To get a more complete view of the Bluetooth communication, I later hooked the app's `onCharacteristicChanged()` function, which logs changes to a Bluetooth Service's Characteristic. This allowed me to track device-side byte payloads and gain a better

understanding of the full communication process between the phone (central) and the device (peripheral).

The first step in decoding the payload was to investigate what was happening between the sends. Frida provides a helpful function for this called `getStackTraceString()`. However, Frida’s documentation can be sparse and often difficult to navigate without prior experience. Fortunately, others have encountered similar issues and created cheat sheets and niche guides for many common Android reverse engineering tasks. Figure 5.11 is an example from an early test. The call stack was generated by triggering an exception, which Frida then intercepted to return the stack values:

Figure 5.11: Payload from Hooked Send Method

This stack trace offered strong clues about which packages to examine for more dispensing information. However, there was still a significant amount of investigation required. For example, the first method visible aside from `send()` was `dispenseFor()`. This is a native method, which means it is implemented in a language other than Java and compiled for a specific computer architecture, usually in C or C++. Native methods can access system-specific APIs not directly available in Java. According to IBM, “Native methods are Java™ methods that start in a language other than Java. Native methods can access system-specific functions and APIs that are not available directly in Java” [9].

These methods are stored in .so files within the APK and cannot be decompiled with Dalvik bytecode decompilers like JADX. Instead, they must be reverse engineered from assembly-level code, which is possible but time-consuming. Because of this, I chose to focus on gathering more context through dynamic analysis rather than attempting to decompile the native code.

To begin understanding the `dispenseFor()` function, I examined its declaration in JADX and logged its parameters. The only reference to it outside the `executeDispense()` function appeared in the BeamLipsController: **public native int dispenseFor(int i, String str, int i2, int i3, int i4, float f, int i5);** This signature provided insight into the types of variables used. The reference within `executeDispense()` offered a clearer picture of the parameters: `beamLipsController.dispenseFor(beamLipsDeviceId, colorUniverse, Color.red(color), Color.green(color), Color.blue(color), floatValue, intValue);` This helped identify the values that would need to be verified during testing. Below is the Frida hook I implemented to intercept calls to `dispenseFor()` and log the relevant information:

```
// Hook dispenseFor(...)

BeamLipsController.dispenseFor.overload(
    'int', 'java.lang.String', 'int', 'int', 'int', 'float', 'int'
).implementation = function (deviceId, colorUniverse, r, g, b, vol, dose) {
    console.log("\n--- dispenseFor(...) called ---");
    console.log(" Device ID: " + deviceId.toString(16));
    console.log(" Color Universe: " + colorUniverse);
    console.log(" Red: " + r + " | Green: " + g + " | Blue: " + b);
    console.log(" Volume (float): " + vol + " → " + toHex(floatToBytes(vol)));
    console.log(" Dose (int): " + dose + " → " + toHex(intToBytes(dose)));
    console.log(" Call Stack (starting with exception):\n" +
        Log.getStackTraceString(Exception.$new()));
}
```

```
logObjectFields(this);  
  
return this.dispenseFor(deviceId, colorUniverse, r, g, b, vol, dose);  
};
```

And below is the resulting output from the test (excluding the call stack for brevity).

```
--- dispenseFor(...) called ---  
Device ID: 431f154e  
Color Universe: cool_nude  
Red: 228 | Green: 101 | Blue: 135  
Volume (float): 0.0799999821186066 -> 3d a3 d7 0a  
Dose (int): 1 -> 00 00 00 01
```

Table 5.1: Cartridge Metadata and Status at Time of Dispense

Field	Cartridge 0	Cartridge 1	Cartridge 2
Type	00	00	00
PCrc16	DBB13410	8E0746A3	258719E9
Volume	4.165000	5.618000	4.084000
Crc16	DD21	CF18	B1CC
DeviceId	431F154E	431F154E	431F154E
TubeId	0	1	2
Last use	2025-06-26 01:29:16	2025-06-26 01:29:16	2025-06-26 01:29:16
Open date	2025-05-20 08:00:00	2025-05-20 08:00:00	2025-05-20 08:00:00
End date	2025-06-27 08:00:00	2025-06-22 08:00:00	2025-06-27 08:00:00
Exp. date	2025-06-27 08:00:00	2025-06-22 08:00:00	2025-06-27 08:00:00
Formula id	MA_513	MA_200	MA_100
Usable Vol.	5.800000	5.800000	5.800000
Duration use	730	730	730
Prod date	2022-06-28 08:00:00	2022-06-23 08:00:00	2022-06-28 08:00:00
Batch id	62W601	62W601	62W601
Color	00DF75A0	006A2623	00C8725C
Present?	1	1	1
Up to date?	1	1	1
Compatible?	1	1	1
Opened?	1	1	1
Expired?	0	1	0
Empty?	0	0	0
Valid?	1	1	1
Virtual?	0	0	0
CanDispense?	1	1	1

It is clear that much of the information required for the dispensing process is stored within the BeamLipsController object. While examining the JADX decompilation, I found that many of these fields also appear in the BeamLipsDevice model:

Figure 5.12: Beam Lips Device Model Jadx

These variables provide valuable clues about what to look for in the raw payload. A crucial detail is that the device information appears to be divided into three cartridges, each maintaining its own set of data. Among the variables that seem particularly relevant to the dispense function are the usable volume, color, and dispense duration. I made note of these values to check whether they could be identified within the payload.

With these values hooked, the next step was to run tests and analyze the outcomes. After altering the variables, I was able to trace how specific values appeared in the payload. For example, the dispense payload associated with the dispenseFor hook revealed several important patterns:

Figure 5.13: Initial Payload from dispenseFor(...) Hook

One of the first observations was that each send() transmission began with the byte aa, followed by a byte that incremented with each new message (e.g., 2e, 2f, 30). This indicates that the first two bytes of each payload likely serve as a counter to keep track of the sequence of transmitted packets. Additionally, after the app sends these two bytes (such as aa 2e), the device responds by writing back the same two bytes, further supporting the idea that this sequence acts as a counter and confirmation mechanism to ensure messages are being received in the correct order.

Another significant detail was that the device ID appeared three times in the second send() payload. Figure 5.15 shows a screenshot of the device's advanced settings screen from the app, which aligns with earlier findings from the hook indicating that the device splits its data into three channels, one for each cartridge. However, I continued investigating to determine whether any additional relevant information was present in the larger send() packet:

Figure 5.14: Payload Two from dispenseFor(...) Hook

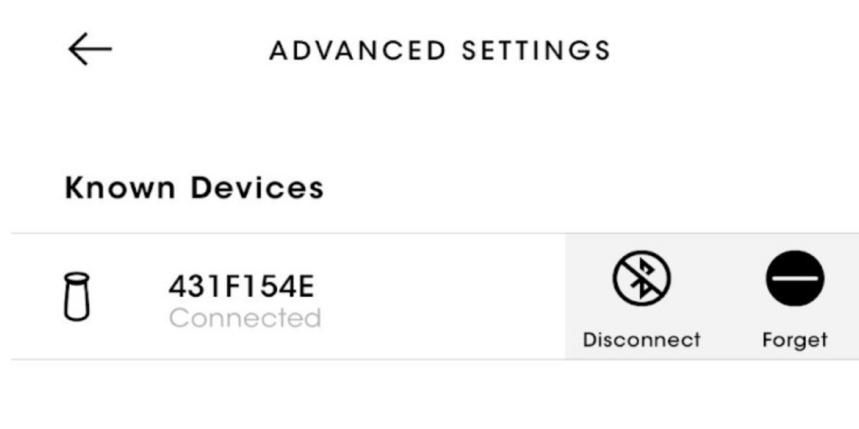


Figure 5.15: Device ID from the App

Working from the hypothesis that the large send payload is organized by cartridge, I continued parsing through the data. As anticipated, I identified the patterns aa 00, aa 01, and aa 02 preceding each occurrence of the device ID, suggesting that the subsequent bytes are associated with data for each cartridge.

It is important to note that, while the information is presented here in a relatively straightforward manner, confirming each hypothesis required numerous tests and careful analysis:

Figure 5.16: Payload Three from dispenseFor(...) Hook

This is where the analysis began to grow more complex. To start, it's important to

understand the expected order in which the data appears. As a reminder, the send() and receive payloads under examination here are part of Bluetooth's GATT profiles. According to the Bluetooth Core Specification 6, the GATT framework “defines procedures and formats of services and their characteristics. The procedures defined include discovering, reading, writing, notifying and indicating characteristics, as well as configuring the broadcast of characteristics” [3]. The specification also states that “Multi-octet fields within the GATT profile shall be sent least significant octet first (little-endian) with the exception of the Characteristic Value field. The Characteristic Value and any fields within it shall be little-endian unless otherwise defined in the specification which defines the characteristic” [3]. This indicates that much of the data should be expected in little-endian format.

Revisiting the payload, the next task was to determine where the dynamic data was stored. After running multiple tests, I identified a pattern in the dispense payloads: the four bytes immediately following 06 00 consistently changed, as did the four bytes preceding 00 00 06. This pattern suggests that these bytes contain data significant to the device and, therefore, crucial to decode if the goal is to replicate the dispense process with a custom payload. When decoding, these bytes can be grouped together, while fixed values can generally be treated as headers or spacers.

Figure 5.17: Payload Four from dispenseFor(...) Hook

Using the cartridge data as a reference for what might appear in the larger payload, and

knowing that volume is one of the most critical aspects of a dispense, I investigated whether the volume could be located within the payload. Because the volume changes between sends for the cartridges involved in dispensing, it was logical that the corresponding data would be part of the bytes that vary after each dispense. During testing, I noticed that when dispensing from only one cartridge, just two bytes would change within one of the four-byte sequences, while the remaining bytes stayed constant across dispenses. With a strong hypothesis about which bytes were significant, the next step was to test various ways these bytes could represent the volume information.

Identifying the correct data took several attempts. One significant obstacle was that, initially, I focused on locating the dispense volume itself, rather than recognizing that the volume appeared as the current total volume within each cartridge. With the realization that I should be looking for information in the cartridge data, representing the total volume per cartridge, and recalled that Android systems generally encode data in little-endian, the discovery became much clearer. Figure 5.18 illustrates the calculation of the volume derived from the payload. These calculations clearly line up with the 'Volume' numbers in

Table 5.1

```
----- BleManager.send(...) called -----
Device ID: 1126110542 (hex: 431f154e)
bytes (hex): aa 2f a5 48 00 00 aa 00 43 1f 15 4e 05 4f 2b 4f b3 85 5d 68 00 00 06 00 f5 0f 63 20 00 00 aa 01 43 1f 15
4e 05 4f 26 4f b3 85 5d 68 00 00 06 00 f2 15 99 c1 00 00 aa 02 43 1f 15 4e 05 4f 2b 4f b3 85 5d 68 00 00 06 00 f4 0f 4d
bf

----- onCharacteristicChanged called -----
Received frames: aa 2f a5 00 0c f5 0f 63 20 f2 15 99 c1 f4 0f 4d bf

----- BleManager.send(...) called -----
Device ID: 1126110542 (hex: 431f154e)
bytes (hex): aa 30 02 00

----- onCharacteristicChanged called -----
Received frames: aa 30 02 00 2c 4c 27 4f 72 c3 a9 61 6c 00 52 53 4d 00 32 32 34 37 30 32 34 35 4c 00 33 00 33 2e 31 31
32 00 32 2e 31 32 00 30 00 44 56 54 2d 32 00

##### - counter
##### - device ID
aa 0* = cartridge number
##### = payload values that change every dispense
##### = remaining cartridge volume

----- Cartridge Little_endian Big_endian Decimal Volume -----
0 f5 0f 0f f5 4085 4.165000 - 0.08 (dispense volume) = 4.085
1 f2 15 15 f2 5618 5.618000
2 f4 0f 0f f4 4084 4.084000
```

Figure 5.18: Payload Five from the dispenseFor(...) Hook

After decoding the volume, there remained only one four-byte sequence and one

two-byte sequence whose purposes were still unknown. To identify these, I revisited the cartridge information for additional clues. One recurring element in the cartridge data was information related to dates. Based on that insight, I experimented with reversing various combinations of bytes into big-endian format and using Python to convert these values into UTC datetime stamps. Ultimately, the full four-byte sequence before 00 00 06 was found to translate to the current print time.

```

Device ID: 1126110542 (hex: 431f154e)
bytes (hex): aa 2f a5 48 00 00 aa 00 43 1f 15 4e 05 4f 2b 4f b3 85 5d 68 00 00 06 00 f5 0f 63 28 00 00 aa 01 43 1f 15 4e 05 4f
26 4f b3 85 5d 68 00 00 06 00 f2 15 99 c1 00 00 aa 02 43 1f 15 4e 05 4f 2b 4f b3 85 5d 68 00 00 06 00 f4 0f 4d bf

----- onCharacteristicChanged called -----
Received frames: aa 2f a5 00 0c f5 0f 63 28 f2 15 99 c1 f4 0f 4d bf

----- BleManager.send(...) called -----
Device ID: 1126110542 (hex: 431f154e)
bytes (hex): aa 30 02 00

----- onCharacteristicChanged called -----
Received frames: aa 30 02 00 2c 4c 27 4f 72 c3 a9 61 6c 00 52 53 4d 00 32 32 34 37 30 32 34 35 4c 00 33 00 33 2e 31 31 32 00
32 2e 31 32 00 30 00 44 56 54 2d 32 00

##### - counter
##### - device ID
aa 0* = cartridge number
##### = payload values that change every dispense
##### = remaining cartridge volume
##### = datetime
|
Cartridge Little_endian Big_endian Decimal Volume
0 f5 0f 0f f5 4085 4.165000 - 0.08 (dispense volume) = 4.085
1 f2 15 15 f2 5618 5.618000
2 f4 0f 0f f4 4084 4.084000

Little_endian Big_endian datetime.utcnow().timestamp()
b3 85 5d 68 68 5d 85 b3 2025-06-26 17:38:59

```

Figure 5.19: Payload Six from the dispenseFor(...) Hook

With this understanding transmission over Bluetooth, it is feasible that if a third party application can be connected to the printer it could replicate these packets of data and communicate with the device. My approach will be to inject code into the native application to modify its behavior.

Despite this shift in strategy, all the insights from analyzing the payloads remain valuable, as they help clarify the types of information the device expects to receive. This data has been documented for possible future use in crafting spoofed payloads. Figure 5.20 shows one of my comprehensive tests analyzing the byte payload from a pink dispense.

Comparing Payloads for CN393 (Pink)Payload 1

```
Raw bytes (hex): aa 0f a5 48 00 00 cartridge0 deviceID 05 4f 2b 4f 47 68 5c 68 00 00 06 00 e9 11 75 5d 00 00 cartridge1 deviceID 05 4f 26 4f 47 68 5c 68 00 00 06 00 f2 15 6e 09 00 00 cartridge2 deviceID 05 4f 2b 4f 47 68 5c 68 00 00 06 00 f4 0f ba 77
```

Payload 2

```
Raw bytes (hex): aa 15 a5 48 00 00 cartridge0 deviceID 05 4f 2b 4f c5 68 5c 68 00 00 06 00 49 11 b5 7a 00 00 cartridge1 deviceID 05 4f 26 4f c5 68 5c 68 00 00 06 00 f2 15 d0 33 00 00 cartridge2 deviceID 05 4f 2b 4f c5 68 5c 68 00 00 06 00 f4 0f 04 4d
```

Payload 3 (volume increase)

```
Raw bytes (hex): aa 18 a5 48 00 00 cartridge0 deviceID 05 4f 2b 4f a0 69 5c 68 00 00 06 00 e5 10 c3 50 00 00 cartridge1 deviceID 05 4f 26 4f a0 69 5c 68 00 00 06 00 f2 15 94 51 00 00 cartridge2 deviceID 05 4f 2b 4f a0 69 5c 68 00 00 06 00 f4 0f 40 2f
```

Payload 4

```
Raw bytes (hex): aa 1b a5 48 00 00 cartridge0 deviceID 05 4f 2b 4f a2 6d 5c 68 ff ff 06 00 95 10 ef 14 00 00 cartridge1 deviceID 05 4f 26 4f a2 6d 5c 68 00 00 06 00 f2 15 f1 13 00 00 cartridge2 deviceID 05 4f 2b 4f a2 6d 5c 68 00 00 06 00 f4 0f 25 6d
```

```
43 1f 15 4e = device ID
aa 00 = cartridge 0           <- pink cartridge, one being changed
aa 01 = cartridge 1
aa 02 = cartridge 2
```

```
Standard Dispense Volume = .08
Increased Dispense Volume = 0.1000000149011612
```

<u>Little_Endian</u>	<u>Big_Endian</u>	<u>Decimal</u>	<u>datetime.utcnowtimestamp()</u>
47 68 5c 68	685c6847	1750886471	2025-06-25 21:21:11
c5 68 5c 68	685c68c5	1750886597	2025-06-25 21:23:17
a0 69 5c 68	685c69a0	1750886816	2025-06-25 21:26:56
a2 6d 5c 68	685c6da2	1750887842	2025-06-25 21:44:02

<u>Cartridge</u>	<u>Little_endian</u>	<u>Big_endian</u>	<u>Decimal</u>	<u>Volume</u>
0	e9 11	11 e9	4585	4.665000 - 0.08 = 4.585 (had one extra dispense 4.585 -.08 = 4.505)
1	f2 15	15 f2	5618	5.618000
2	f4 0f	0f f4	4084	4.084000
0	49 11	11 49	4425	4.505000 - 0.08 = 4.425
1	f2 15	15 f2	5618	5.618000
2	f4 0f	0f f4	4084	4.084000
0	e5 10	10 e5	4325	4.425000 - 0.1000000149011612 = 4.325
1	f2 15	15 f2	5618	5.618000
2	f4 0f	0f f4	4084	4.084000
0	95 10	10 95	4245	4.325000
1	f2 15	15 f2	5618	5.618000
2	f4 0f	0f f4	4084	4.084000

Figure 5.20: Multiple Payload Analysis

5.6 EXECUTING THE DISPENSE

With the comprehensive data gathered from earlier experiments, I set out to modify the color parameters and initiate a dispense. I once again used Frida's straightforward Java method hooks, but in this iteration I directly manipulated the input variables of the `dispenseFor()` function. For clarity, the method signature is:

```

beamLipsController.dispenseFor(
    beamLipsDeviceId,
    colorUniverse,
    Color.red(color),
    Color.green(color),
    Color.blue(color),
    floatValue,
    intValue
);

```

This function receives the cartridge universe identifier, the red, green, and blue components of the selected shade, the volume to dispense, and the dose count. My objective is to demonstrate the capacity to override these values at runtime. To achieve this, I redirected a dispense request originally set to light pink so that it produced a light brown output instead. Figure 5.21 depicts the test arrangement, highlighting device connections and the sequence of operations.

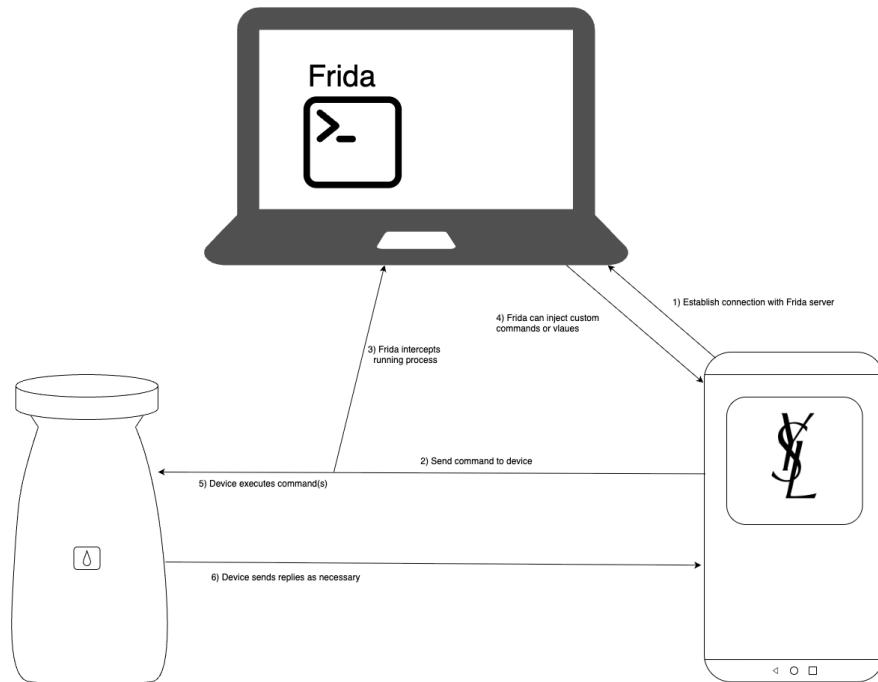


Figure 5.21: Dispense Environment Setup

The Frida script follows the same pattern as previous hooks. First, I capture and log each incoming parameter to verify the original settings. I then replace the RGB values with those corresponding to the CN24 light brown shade, which are 200, 114, and 92 based on an earlier hook session. Finally, I invoke the original method with the modified values. The full script is presented below:

```
Java.perform(function () {
    var BeamLipsController = Java.use("com.vinsol.loreal.
        PersoLips.utils.BeamLipsController");
    var overload = BeamLipsController.dispenseFor.overload(
        'int',
        'java.lang.String',
        'int',
        'int',
        'int',
        'float',
        'int'
    );

    overload.implementation = function(deviceId, colorUniverse,
        red, green, blue, volume, dose) {
        console.log("\n==== dispenseFor(...) called ====");
        console.log("Device ID: " + deviceId);
        console.log("Color Universe: " + colorUniverse);
        console.log("Original Red: " + red);
        console.log("Original green: " + green);
        console.log("Original Blue: " + blue);
        console.log("Original volume: " + volume);
    }
});
```

```
    console.log("Original dose: " + dose);

    console.log("----- Altering Color Values
to print CN24 -----");

    red = 200;
    green = 114;
    blue = 92;

    console.log("Modified Red: " + red);
    console.log("Modified Green: " + green);
    console.log("Modified Blue: " + blue);

    return overload.call(
        this,
        deviceId,
        colorUniverse,
        red,
        green,
        blue,
        volume,
        dose
    );
};

console.log("*** Hooked correct overload of dispenseFor in
BeamLipsController ***");
```

```
});
```

The result was incredibly exciting. As shown in Figure 5.22, the device produced a light brown swatch even though the app remained set to light pink. However, the original light pink also appeared alongside the brown, despite only a single invocation of the modified `dispenseFor()` call. Furthermore, the final shade reported was CN63—a brown-pink hybrid that does not exist on the official color wheel (see Figure 5.24).

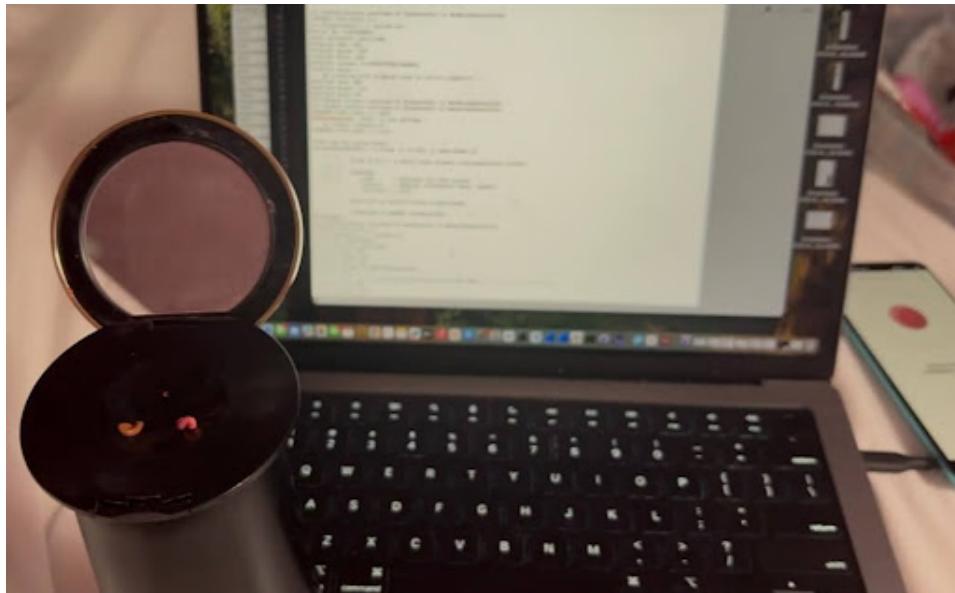


Figure 5.22: Picture of Altered Dispense Output

```
npargas@MacBookPro ~ % frida -U -p 2411 -l swap_shade.js
    /----| Frida 17.0.4 - A world-class dynamic instrumentation toolkit
    | ( )| Commands:
    /|-|_ help      -> Displays the help system
    . . . object?   -> Display information about 'object'
    . . . exit/quit -> Exit
    . . . More info at https://frida.re/docs/home/
    . . . Connected to KB2005 (id=5a4bc47e)
Attaching...
*** Hooked correct overload of dispenseFor in BeamLipsController ***
[KB2005:PID:2411 ]->
== dispenseFor(...) called ==
Device ID: 1126110542
Color Universe: cool_nude
Original Red: 228
Original green: 101
Original Blue: 135
Original volume: 0.07999999821186066
Original dose: 1
----- Altering Color Values to print CN24 -----
Modified Red: 208
Modified Green: 114
Modified Blue: 92
[KB2005:PID:2411 ]-> exit
Thank you for using Frida!
```

Figure 5.23: Frida Console Logs During Altered Dispense

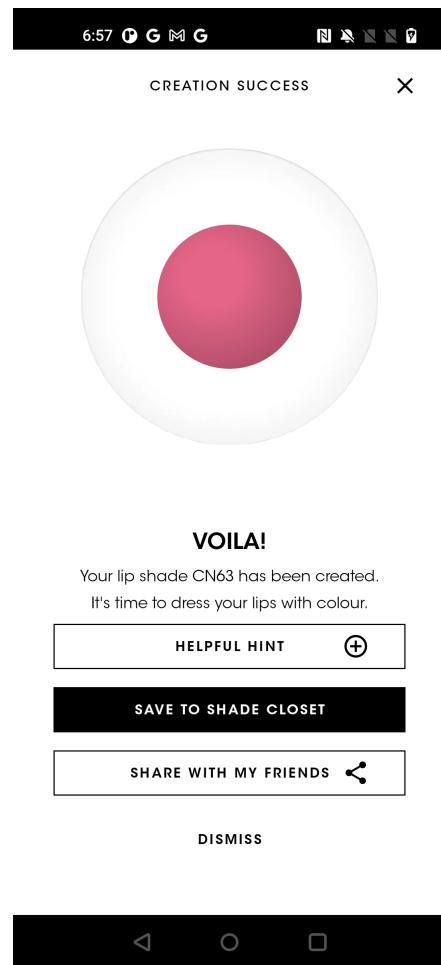


Figure 5.24: cn63 Screen Following Altered Dispense

These findings indicate that additional processing occurs after the RGB parameters are accepted. Possible explanations include internal blending algorithms, use of cached color values, or fallback logic within the application or firmware on the printer itself. To uncover the precise source of these effects, I could trace subsequent method calls and capture post-dispense data packets in future experiments

CHAPTER 6

CONCLUSIONS AND FUTURE WORK

6.1 REVERSE ENGINEERING OUTCOME

This project has undertaken a comprehensive investigation of reverse engineering practices as applied to complex Bluetooth-enabled systems. By examining both theoretical foundations and practical implementations, I have illuminated the interplay between static code analysis, dynamic instrumentation, and wireless packet inspection. The objective was twofold: to understand the internal architecture of a commercial Android application that communicates with a Bluetooth peripheral and to demonstrate the feasibility of intercepting and manipulating its behavior in real time.

First, static analysis using JADX provided a crucial starting point for mapping the application's data flow and identifying key functions. Through systematic decompilation and code review, I traced how user inputs and device commands are translated into Bluetooth payloads. This phase revealed the structure of the proprietary protocol and established a blueprint for subsequent dynamic testing. However, static techniques alone cannot capture runtime behaviors such as input validation, encryption routines, or environment-specific logic.

To address these limitations, I employed Frida for dynamic instrumentation. By injecting hooks into the application's Java methods, I was able to observe live function calls,

inspect variable values, and modify parameters before transmission. Frida’s versatility enabled close inspection of runtime behavior and the execution of targeted manipulations, confirming that the application’s control logic could be influenced without source-level modifications. This dynamic approach complemented the insights gained from static decompilation and provided definitive evidence of our capacity to alter an application’s operation.

Complementing code-centric methods, Bluetooth packet sniffing yielded an unambiguous record of over-the-air exchanges between the mobile application and the peripheral device. Using a hardware sniffer and analysis software, I captured raw packets that corroborated the structure inferred from code review and Frida instrumentation. These traces ensured that our reverse-engineering conclusions were grounded in empirical measurements rather than extrapolation. The integration of packet-level data with code-level insights formed a robust, end-to-end understanding of the entire communication pipeline.

Throughout this work, an iterative methodology proved essential. Each new discovery prompted reevaluation of prior assumptions and guided subsequent experiments. Establishing a reliable environment—including rooted devices and virtual machines—facilitated rapid prototyping and reproducibility. Equally important was meticulous documentation of all procedures, findings, and hypotheses, which preserved critical knowledge and supported collaborative scrutiny.

In sum, this project demonstrates that commercial-grade Bluetooth applications can be effectively reverse engineered through a judicious combination of static analysis, dynamic instrumentation, and packet monitoring. The techniques and workflows developed here provide a foundation for further research into more advanced protocols and emerging hardware platforms. With refined automation tools and expanded coverage to additional communication layers, future work can continue to push the boundaries of applied reverse engineering within the security research community.

6.2 FUTURE WORK

The current study has demonstrated that manipulating the dispense process of the device is not only possible but can serve as a robust foundation for achieving comprehensive command over its functions. Building upon these initial findings, the next phase of research could refine communication protocols, reverse-engineer proprietary algorithms, develop a standalone application, and analyze firmware to uncover hidden operational logic.

First, building on the initial packet capture efforts conducted during this study, the research could return to in-depth Bluetooth packet analysis to further elucidate the device's communication protocol. This effort might involve capturing and parsing additional raw GATT traffic to confirm previously observed services, characteristics, and command sequences. Using tools such as Wireshark or specialized BLE sniffer, the study could refine documentation of packet structures, attribute handles, and payload formats. Second, additional efforts could concentrate on isolating and interpreting the device's color-related commands within the captured Bluetooth packets. This would involve identifying specific GATT characteristics or payload patterns that correspond to shade adjustments, enabling direct control over pigment ratios without requiring extensive algorithmic reconstruction.

Third, with both communication protocols and color models established, the research could focus on developing a standalone mobile application independent of the vendor's software. This application could include detailed documentation of the Bluetooth services and characteristics, such as UUIDs, properties, and data formats, required to interface with the device. The user interface could then offer intuitive workflows for shade selection, custom color creation, and real-time monitoring. Finally, to reveal any embedded processing routines not exposed through the GATT interface, the study could perform firmware extraction and static analysis. Through hardware teardown and debugging techniques, the research could retrieve the device's firmware image.

Reverse-engineering tools could then be employed to identify internal state machines, calibration sequences, and safety checks. Mapping these routines to observed device

behaviors could close gaps in the external command model and inform enhancements to custom control.

In conclusion, the proposed future work lays out a clear path for turning early experiments into a fully independent, user-controlled system for accurate color dispensing. By improving communication with the device, learning how its internal processes work, creating a custom app, and examining its internal software, this research could move entirely away from relying on the original manufacturer's app and open up new possibilities for personalized cosmetic production.

REFERENCES

- [1] Mohammad Afaneh, Kevin Townsend, and M. Karvinen. *Building Bluetooth Low Energy Systems*. Packt Publishing, 2017. ISBN: 9781786461087. URL: <https://learning.oreilly.com/library/view/building-bluetooth-low/9781786461087/> (pages 33, 54).
- [2] Bluetooth SIG. *Bluetooth Core Specification Version 6.0*. Accessed: 2025-04-24. 2024. URL: <https://www.bluetooth.com/specifications/specs/core60-html/> (pages 32, 37).
- [3] Bluetooth SIG, Inc. *Core Specification 5.4 – Host*. Technical Report. Available online. Bluetooth Special Interest Group, Feb. 2023. URL: <https://www.bluetooth.com/wp-content/uploads/Files/Specification/HTML/Core-54/out/en/host.html> (page 97).
- [4] Bruce Dang et al. *Practical Reverse Engineering X86, X64, ARM, Windows Kernel, Reversing Tools, and Obfuscation*. John Wiley and Sons, Incorporated, 2014 (pages 11, 18).
- [5] Spiros Daskalakis. *Introduction to Bluetooth Low Energy v1.1*. Accessed: 2025-04-24. 2019. URL: <https://daskalakispiros.com/files/Ebooks/Intro+to+Bluetooth+Low+Energy+v1.1.pdf> (pages 44–47, 50–53).
- [6] Julien Duchêne et al. “State of the art of network protocol reverse engineering tools”. In: *Journal of Computer Virology and Hacking Techniques* 14.1 (2018), pp. 53–68. doi:

- 10.1007/s11416-016-0289-8. URL: <https://doi.org/10.1007/s11416-016-0289-8> (page 14).
- [7] Eldad Eilam. *Reversing : Secrets of Reverse Engineering*. John Wiley and Sons, Incorporated, 2005 (pages 5, 7, 10–11, 16–18, 20–22).
- [8] Dawn Griffiths and David Griffiths. *Head First Android Development*. 3rd ed. Sebastopol, CA: O'Reilly Media, Incorporated, 2021. ISBN: 978-1492076513 (page 57).
- [9] IBM Corporation. *Native methods and the Java Native Interface*. IBM i 7.5.0. Accessed July 26, 2025. IBM. 2025 (page 90).
- [10] Intel. *Intel® 64 and IA-32 Architectures Software Developer's Manual*. Intel, 2024 (page 21).
- [11] Intel Corporation. *How Does Bluetooth Work?* Accessed: 2025-04-24. 2023. URL: <https://www.intel.com/content/www/us/en/products/docs/wireless/how-does-bluetooth-work.html> (page 29).
- [12] Alex Kalinovsky. *Covert Java: Techniques for Decompiling, Patching, and Reverse Engineering*. Indianapolis, IN, USA: Sams Publishing, 2004. ISBN: 0-672-32638-8 (page 55).
- [13] Tim Lindholm et al. *The Java® Virtual Machine Specification, Java SE 8 Edition*. Published electronically by Oracle on 2015-02-13. Boston, MA: Addison-Wesley Professional (Oracle), Pearson Education, May 2014. ISBN: 978-0133922721 (page 55).
- [14] Sasha Sirotkin, Kassem Fawaz, and Anthony Rowe. “Bluetooth 5 and Beyond”. In: *Next Generation of Internet of Things*. ISTE Press - Elsevier, 2020. ISBN: 9781789451771. URL: <https://learning.oreilly.com/library/view/next-generation-of/9781789451771/c06.xhtml#s6-5> (pages 31–33, 43, 48).
- [15] Kevin Townsend et al. *Getting Started with Bluetooth Low Energy: Tools and Techniques for Low-Power Networking*. O'Reilly Media, Inc., 2015. ISBN: 9781491900550. URL: <https://learning.oreilly.com/library/view/getting-started-with/9781491900550/> (pages 38–39, 49, 52).

- [16] Reginald Wong. *Mastering Reverse Engineering*. Packt Publishing, 2018 (pages 12, 20).

